ALGORITHMIC MUSIC COMPOSITION

by:

S. Jacholke   23671750

submitted in persuit of the degree

BACHELOR OF ENGINEERING
in
COMPUTER AND ELECTRONIC ENGINEERING

North-West University Potchefstroom Campus

Supervisor: Mr. A.J. Alberts
Potchefstroom
October 25, 2015
version 2.2

Where words fail, music speaks.

— Hans Christian Andersen

DECLARATION

I, Stefan Jacholke hereby declare that the report all work associated with the 'Algorithmic Music Composition' project are my own original work and had not already been submitted to any university or institution for examination

Potchefstroom, October 25, 2015

Stefan Jacholke

ABSTRACT

Obtaining the required music licensing may not be viable for certain individuals and business owners. Alternatively the music can be composed and produced, however this requires time and skill.

For this project another alternative is investigated, to produce music using computer algorithms.

The field of algorithmic music composition is not new. The problem of algorithmically composing music is difficult and there is a wide range of research on the topic.

For this project, the relevent research is investigated. Possible algorithms were identified and implemented. Algorithms producing pleasant results include: Markov Models, a simple feed forward neural network, Markov Chains and Genetic Algorithms. These algorithms were kept and implemented in a end-user application.

The Windows application developed features a friendly user interface that allows the user to compose, playback and store simple monophonic melodies in a certain style. In order to optimize the end user experience the algorithms were pretrained were possible.

Since it is difficult to quantify the pleasantness of a song, due to the inherent subjective experience, it was necessary to rely on pattern recognition and machine learning techniques. A weakness of the algorithms implemented is that they do not have a overarching theme or structure.

CONTENTS

LIST OF FIGURES

LIST OF TABLES

LISTINGS

LISTE DER ALGORITHMEN

ACRONYMS

ANN Artificial Neural Network

GP Genetic Program

GA Genetic Algorithm

HMM Hidden Markov Model

MIDI Musical Instrument Digital Interface

NCD Normalized Compression Distance

NID Normalized Information Distance

IGA Interactive Genetic Algorithm

SRN Simple Recurrent Network

LSTM Long Short Term Memory

ART Adapative Resonance Theory

RNN Recurrent Neural Network

XAML Extensible Markup Language

WPF Windows Presentation Foundation

XML Extensible Markup Language

UI User Interface

LOC Lines of Code

FFANN Feed Forward Neural Network

Part I

INTRODUCTION

ALGORITHMIC MUSIC COMPOSITION

1.1 INTRODUCTION

Music is the art form of sound. Certain characteristics and patterns can be identified that belong to certain types of music.

There has been a lot of work done in computer generated music and computer assisted music composition.The design of computational methods for musical style imitation has been far more difficult than initially imagined. Recent research into music composition by means of evolutionary genetic algorithms and other machine learning algorithms such as neural networks have met some success.

Music generated algorithmically by computers might some day share the success known by modern and historic music. In a hypothetical time period artists might compete with algorithmic composers.

Computer generated music is an old idea, however existing solutions that generate music use outdated procedural algorithms and do not match the quality of proper historic music. The variety and quality of existing solutions is limited. Most current solutions were implemented to accommodate the research being done and are not suitable for end-users.

This section will outline the problems that are encountered with conventional music composition and a solution is proposed. Some of the ideas behind systems that learn music composition will be outlined and described. The problem of implementing and creating a solution will be detailed as well as the steps required to complete such a project.

1.2 PROBLEM

In order to avoid copyright infringement music has to be licensed for use in commercial applications. Small business owners, independent developers and other smaller organizations may not have the required funding in order to obtain the relevant licensing.

The alternative solution is to produce the required music self, however this requires time and skill.

A solution to this problem is to generate music by artificial intelligent means. Algorithmic music composition is of theoretical interest and research has been done for composing music by genetic algorithms and neural networks, amongst other algorithms.

A software application that is able to compose and play monophonic melodies according to a certain style by means of machine learning algorithms could solve the licensing and resource problem.

Figure 1: Figure of high-level architecture of project

1.3 SOLUTION

The solution to the problem is to create an application that is capable of composing and playing simple monophonic melodies using machine learning algorithms. Algorithms other than learning algorithms may be used, however these algorithms are narrow in application. Using learning algorithms and pattern recognition allows the application to produce melodies according to a selected style.

There is a need for a user-friendly application that can be used to generate simple monophonic melodies according to a certain style or category of music. The application is required to have the following features :

1. Melodies are to be algorithmically composed according to a certain style.
2. Provide playback functionality as well as saving and loading of previously composed songs.
3. The user should be able to select the style
4. The application should utilize a MIDI library which will be used as training data for the machine learning algorithms.

Figure 1 indicates the interaction between the user and the software.

Since the quality and aesthetics of a musical melody is subjective, it is beyond the scope of the application to ensure that each melody produced is objectively good[1]. Composed musical pieces will take the form of simple melodies.

1.4 ALTERNATIVE SOLUTIONS

A wide variety of research has been done and in this section we will briefly look at some of the existing solutions and their drawbacks

---

1 The focus is rather on the generated musical pieces not being objectively bad

1.4.0.1 *NEUROGEN*

NEUROGEN is a system that was designed to compose small diatonic, western-type, four part harmony compositions based on a training set of example musical fragments. NEUROGEN tries to produce coherent music of the type that is typically found in traditionally hymns [2]. However NEUROGEN is only tries to fulfill its main research goals.

1.4.0.2 *CONCERT*

CONCERT incorporates psychologically-grounded representations of pitch, duration and harmonic structure in order to compose novel melodies. CONCERT uses a neural network in order to do note by note predictions.

CONCERT struggles to compose natural music as the music produced was not globally coherent [3].

CONCERT also only tries to fulfill its research goals, that is, to compose music on a note by note basis using a neural network.

More recent studies into recurrent neural networks based on LSTM indicate that it is indeed possible to compose music using recurrent neural networks that have global structure [1].

1.4.0.3 *GenJam*

GenJam - Genetic Jammer is a IGA that learns to improvise jazz [4]. The system is able to:

1. Play full-chorus improvised solos
2. Respond interactively when the trumpet is played to trade fours or eights
3. Perform a smart echo of improvisation

The main drawback of GenJam is that it is an interactive algorithm and as such requires input from the user. The creates a performance bottleneck as it is time consuming for the user to rate musical pieces. See section 5.6 for work on IGAs

1.4.1 *jMusic*

jMusic is an application and library that facilitates musical composition in Java [5]. The application is designed to provide composers and developers a plethora of compositional tools that can be used for generative music, instrument building, interactive performance and music analysis. In addition the software provides the following:

- Familar music data structures
- Organizing, manipulating and analyzing music data
- Real time playback of music scores
- Read, write MIDI, Extensible Markup Language (XML) and .jm files

jMusic is mostly a developer tool to facilitate composition. It does not feature a user interface which can be used to easily generate compositions. Furthermore since it is mostly a tool to aid music composition it does not employ machine learning algorithms. The library requires exact and descriptive input in order to generate melodies, and as such requires domain knowledge.

1.4.1.1 *Other*

Most of the work done in the field has been aimed on narrow application of the ideas being investigated. Thus algorithms and applications constructed were aimed only to experimentally study the ideas being investigated and not to serve as a end-user product.

Some other small application exist which employ Interactive Genetic Algorithms (IGAs) such as:

1. Evolutune[2]
2. Song Builder[3]
3. DarwinTunes[4]

Even though more recent research into algorithmic music composition yields promising results there is no single application combining the more recent research into a coherent application that is able to play monophonic melodies[5].

1.5 FEASIBILITY

There has been a plethora of research done into algorithmic music composition. Some of the more recent research that utilizes machine learning algorithms is investigated in chapter 13.

Some machine learning methods to compose music algorithmically include:

- Recurrent neural networks [1, 3, 6]
- Genetic algorithms using NCD as fitness functions [7, 8]
- Genetic algorithms using neural networks as fitness functions [2, 9, 10]
- Genetic algorithms employing Zipf's law and cosine similarity [8, 11, 12]
- Interactive genetic algorithms [4, 9, 13, 14, 15]
- Markov Chains for music composition [16, 17].
- Hidden Markov Models for harmonization [18].
- LSTM neural network for composing blues [1].

---

2 http://askory.phratry.net/projects/evolutune/
3 http://www.compose-music.com/
4 http://game.darwintunes.org/
5 In addition there is no accessible user-friendly application that expose machine learning based algorithmic composition to the user

A variety of ideas and techniques in algorithmic music composition have already been researched. The difficulty arises in the implementation of the ideas (as the process is only fundamentally documented), representation of music in data structures and the implementation of the relevant machine learning algorithms.

1.6 METHODOLOGY

In Royce's classical waterfall model the following steps are followed when developing software:

- Requirements specification
- Design resulting in the software architecture
- Construction
- Integration
- Testing and debugging
- Installation
- Maintenance

Since there are a variety of algorithmic music composition strategies, and no formal analysis that compares the quality of the resulting music of each strategy (difficult as it is a subjective matter) the software prototyping methodology could be used to test these different ideas and to construct a trade-off study to determine the most feasible techniques.

Software prototyping is the development of prototypes - incomplete versions of the software. Some basic principles of software prototyping include:

- Prototype selected parts
- Breaking project into smaller segments
- User or client may be involved
- Iterative modification to meet user demands

The main ideas of the systems engineering process, as applicable to software will be followed. Ideas from the waterfall model and software prototyping will be incorporated.

The basic formulation of the design of the project is as follows:

1. Conceptual design
    - Identify problem
    - Identify requirements
    - Resource allocation
    - Literature study
    - Feasibility analysis
    - Trade off study

2. Preliminary design
3. Detail design
   - Design of software architecture
   - Prototyping of different strategies and algorithms
4. Construction
   - Code is written
   - Code is tested
   - Iteratively improve until project meets specifications
5. Phase out, support, maintenance

In order to identify suitable algorithms a literature review will be done. The various algorithms will be prototyped and tested and the best few resulting algorithms will be used. Software development will follow the waterfall model.

1.7 CONCLUSION

Since the act of generating music using algorithmic means is of theoretical interest there is a plethora of research available, however most of these implementations only aim to explore the main idea being researched. Thus there is a lack of accessible, user-friendly applications that are able to generate music algorithmically.

The aim of this project is to construct an application that is able to compose music algorithmically. The application will utilize machine learning algorithms in order to compose monophonic melodies according to the style of a certain author or theme. These ideas have already been researched and thus the application is feasible. The only cost is time and effort.

The methodology for successfully completing the project is outlined and a schedule for completing each goal is presented. This should assist in meeting the deadlines and completing the project successfully according to the specifications.

Part II

LITERATURE REVIEW

2

MUSIC AND MUSIC-REPRESENTATION

2.1 MUSICAL ELEMENTS

Music is an art form for which the medium is sound. The common elements in music are:

- Pitch - Subjective sensation reflecting lowness and highness of sound, also represented more objectively as frequency.
- Rhythm - Arrangement of sounds and silences.
- Dynamics - Execution of a given piece (speed, volume)
- Timbre - Tone

Music composition refers to the creation and recording of music through a medium that others can interpret. Music can be composed for repeated playback or it can be composed on the spot.

A melody is a set of notes (or rests) that are performed in series. Each note may have a different length and different stress. These notes are arranged in a certain rhythmic pattern.

Notes were traditionally given a letter to represent the pitch of the note. The names come from the set $\{A, B, C, D, E, F\}$. Notes may have their pitch modified by additional symbols such as a sharp(#).

Notes and rests may have different lengths.

Li and Sleep have found that a given piece of music remains recognizable when the length of the notes are randomized [19].

2.2 MIDI

Music can be stored digitally in a variety of formats and encodings. This allows repeated performance of the same track and also eases distribution of the music.

MIDI is a standard which describes the protocol, digital interface and connectors that allows electronic instruments to communicate with one another.

The MIDI 1.0 detailed specification fully describes the MIDI interface [20]

The MIDI file format describes the way MIDI information is stored in a file. Each MIDI file starts with a header chunk that describes the time division and the number of tracks. After the header chunk multiple track chunks occur. Each track contains multiple MIDI events.

The header chunk is described as follows:

```
"MThd" + [header_length] + [format] + [n] + [division]
```

where

[header_length] always 6 bytes

[format] 0 - single track format, 1 multiple track format, 2 multiple song format

[n] number of track chunks

[division] unit of time for delta timing

After the header chunk [n] track chunks follow. Each track chunk is composed as follows:

```
"MTrk" + [length] + [track_event] <+ [track_event1] + [track_event2] +
    [...] + [track_eventn]>
```

where

[length] number of bytes in track chunk

[track_event] sequenced track event, described next

The track chunk can be seen to contain multiple track events. Each track event consists a delta time and either a midi event, meta event or sysex_event:

```
[v_time] + [midi_event] | [meta_event] | <sysex_event>
```

The meta-event contains additional information such as text which can be displayed, instrument names, lyrics and so on.

Table 1: Table containing some midi events

Command	Meaning	parameters	param 1	param 2
0x80	Note-off	2	key	velocity
0x90	Note-on	2	key	velocity
0xA0	Aftertouch	2	key	touch
0xB0	Continuous controller	2	controller #	controller value
0xC0	Patch change	2	instrument #	
0xD0	Channel Pressure	1	pressure	

Table 1 shows some of the possible midi events. Each event has a command and additional arguments. Note-on and note-off are the two more important events in the table and indicate how long a note is played.

The full midi specification can be ordered online at midi.org[1]

2.3 REPRESENTATION

Representation of music is important to achieve successful generation of a correct solution [2]. However the search space of machine learning algorithms may be too large if limitations are not introduced [21].

---

[1] http://www.midi.org/

Figure 2: Excerpt of Bach's BWV 05225

Figure 2 shows the music sheet of an excerpt of Bach's BWV0525 Sonata using modern notation.

The musical data from musical pieces need to have a proper representation in order to be used by a machine learning algorithm.

In [1] Eck a form of 12-bar blues was used with 8 notes per bar. He used the following representation for the **LTSM!** system:

1. 12-bar musical pieces were used
2. 8 notes per bar were used
3. The same chords were used between songs, see figure 5
4. The melody notes were built using the pentatonic scale
5. The neural network inputs indicate whether a certain note is on or off

A melody would be represented differently in a linear genetic algorithm. A simple melody consisting only of a few notes could be represented as seen in figure 3. Note the duration of a note is not defined and could be consider fixed.

However it might be preferable to encode more events such rests, chords, dynamics, note duration and so on.

Tree based representations are found in genetic programming [22]. Minsky has argued that tree representation is more suited to music as it mimics the hierarchical structure that is found in music [23].

A melody could be represented in tree form as seen in figure 4. The bottom leaves denote the pitch and duration. The higher level nodes perform operations.

Johanson and Poli employ several different operators including mirroring, concatenation, repetition and mirroring [15]

Tree based representations do not have a fixed size, as is commonly found in linear genetic representations. Care must be taken to not let the structure become too large. This is typically done by specifying the maximum depth of the tree.

Figure 3: Simple melody represented as a bit string
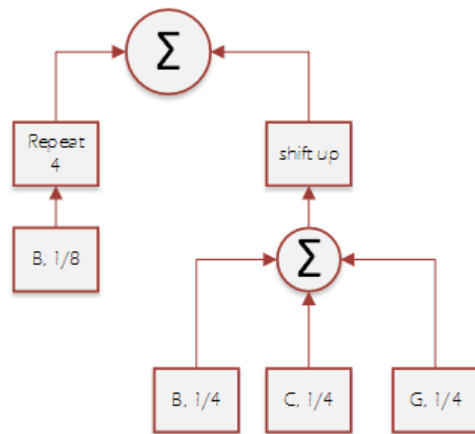
Figure 4: Simple melody represented in tree form

Figure 5: Training chords used in [1]

MUSIC COMPOSITION ALGORITHMS

3.1 HIDDEN MARKOV MODELS

Figure 6: Example of a Hidden Markov Model

3.1.1 *Introduction*

HMM are stochastic methods that are used to model sequential data such as speech and gesture recognition. Both the underlying process which is hidden is stochastic and the observable process. Only a first order hidden markov model will be considered. Under the Markov property the next state is only dependent on the current state of the system. Such states may not be known and may be hidden from the observer as only the output values are observable. When a event is generated from a state, the model moves into a new state based on its transition probabilities. The term hidden is commonly used to indicate that many different state sequences can generate the same observed sequence of events.

The HMM is formally defined by the triple:

$$\lambda = (A, B, \pi)$$

where A is the transition matrix, B the emission matrix and π the initial state vector. Let S be the set of states, V the set of observations, Q a fixed state sequence of length T and O a corresponding observation sequence. The transition probabilities are then given by:

$$A = P(q_t = s_j | q_{t-1} = s_i)$$

and the emission probabilities are given by:

$$B = P(o_t = v_k | q_t = s_i)$$

and the initial state probabilities by:

$$\pi = P(q_1 = s_i)$$

Given a sequence of observations we want to compute the probability of an observation sequence given the model.

The probability of the observation sequence O for a state sequence is given by:

$$P(O|Q, \lambda) = \prod_{t=1}^{T} P(o_t | q_t, \lambda)$$

The probability of the state sequence is simply a product of the state path:

$$P(Q|\lambda) = \pi_{q_1} a_{q_1 q_2} \cdots a_{q_{T-1} q_T}$$

Thus the probability of the observations given the model can be calculated by:

$$P(O|\lambda) = \sum_Q P(O|Q, \lambda) P(Q|\lambda)$$

The forward algorithm calculates this efficiently by caching previous calculations.

3.1.2 *Training*

A variety of learning algorithms exist which compute the structure of the model and also calculate the emission and transmission matrices. The Baum-Welch is a unsupervised learning algorithm which re-estimates the parameters (A, B, π). The training problem can be solved in terms of joint events and state variables [24].

The joint variable $\xi_t(i, j) = P(q_t = S_i, q_{t+1} = S_j | S, \lambda)$ is given by:

$$\frac{\alpha_t(j) a_{ij} b_j(O_{t+1}) \beta_{t+1}(j)}{\sum_{k=1}^{N} \sum_{l=1}^{N} \alpha_t(k) a_{kl} b_l(O_{t+1}) \beta_{t+1}(l)}$$

The state variable $\gamma_t(i) = P(q_t = S_i | O, \lambda)$ is given by:

$$\gamma_t(i) = \frac{\alpha_t(i) \beta_t(i)}{\sum_{j=1}^{N} \alpha_t(j) \beta_t(j)}$$

Upon which the parameters can be updated. The updated state transition probability is then given by:

$$a'_{ij} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)}$$

Figure 7: Example of a Markov Chain

The updated emission probability is given by:

$$b_j'(k) = \frac{\sum_{t=1, o_t=v_k}^{T} \gamma_t(j)}{\sum_{t=1}^{T} \gamma_t(j)}$$

The updated initial probability by:

$$\pi_i' = \gamma_1(i)$$

3.1.3 *Composition*

HMM have been applied with some degree of success to music composition. A system has been developed for producing a counterpoint line to a cantus firmis in the style of Palestrina [17]. Another approach utilizes a HMM for chorale melody harmonization [18].

3.2 MARKOV CHAINS

3.2.1 *Introduction*

A Markov Chain is a stochastic model describing a sequence of events. The Markov Chain has the property that the probability of the next state depends only on the current state of the chain.

$$Pr(X_n|X_1, X_2, \ldots, X_{n-1}) = P(X_n|X_{n-1})$$

Figure 7 is an example of a Markov Chain. For each state there is a probability of proceeding to the next state. For example if the current state is F, there is a 10% chance of the next state being F again, and a 90% chance of the next state being A.

A Markov Chain can have a order. For a simple note such as figure 7 if the current note is B then the note A has a 30% chance of being next. This is a first order Markov Chain. A more complex chain of order 2 could have the current state being composed of two notes. Thus the next note is dependent on the previous two notes. A N-th order Markov chain can be represented by a transition matrix, which corresponds to a N+1 dimension probability table.

For a Markov Chain of order m the following holds:

$$\Pr(X_n = x_n \mid X_{n-1} = x_{n-1}, X_{n-2} = x_{n-2}, \ldots, X_1 = x_1)$$

$$= \Pr(X_n = x_n \mid X_{n-1} = x_{n-1}, X_{n-2} = x_{n-2}, \ldots, X_{n-m} = x_{n-m})$$

For $n > m$

3.2.2 *Composition*

Markov Chains must be derived from existing music and generate output of the same style as the input. This reduces the compositional value and novelty however it suits the purpose of generating music in a specific style [25].

Context models such as Markov Chains have several advantages that are useful in algorithmic music composition [26]:

- Events can be predicted from preceding events.
- Context models are fast and efficient data structures can be used.
- Most context models are straightforward to generate new music with.

3.3 NEURAL NETWORKS

3.3.1 *Background*

Figure 8: Model of artificial neuron

Figure 9: Figure of ANN in feed-forward topology

An artificial neural networks consists of layers of artificial neurons[1] (see figure 8) that are connected to each other in a certain topology [27].

Figure 9 shows a multilayer ANN in feed-forward topology. The first layer is referred to as the input layer, the right-most layer is referred to as the output layer, the layers in between the input layer and the output layer are known as hidden layers.

Neurons have a number of inputs which are summed into an activation function. The activation function provides the output of a neuron. In figure 8 we have

$$o_j = \varphi \left(\sum_{i=1}^{n} x_i w_{ij} \right)$$

A common activation function is the continuous sigmoid function given by

$$\varphi(t) = \sigma(t) = \frac{1}{1 + e^{-\beta t}} \tag{1}$$

where β is the slope parameter. The derivative of the sigmoid function is easy to obtain and is given by:

$$\frac{\sigma(t)}{dt} = \sigma(t)(1 - \sigma(t)) \tag{2}$$

for $\beta = 1$

In order to determine weights for a two-layer neural network we need to minimize the error between the output of the neural network and the given target value for a set of inputs.

Thus for the output neurons we have

$$E(\vec{w}) = \frac{1}{2} \sum_{d \in D} (t_d - \sigma_d)^2 \tag{3}$$

where D is the set of training examples and t_d is the target output for training example d and o_d is the output of the artificial neuron for training example d We need to minimize equation 3. Obtaining the derivative of E with respect to w_i yields:

$$\frac{\partial E}{\partial w_i} = \sum_{d \in D} (t_d - o_d)(-x_i d)$$

---

1 neurons are commonly referred to as units

where x_{id} is the input for component x_i

Gradient descent is a optimization algorithm that takes steps proportional to the negative of the gradient. This we update the weights by:

$$\vec{w} \leftarrow \vec{w} - \eta \nabla E(\vec{w})$$

Applying the gradient descent algorithm we obtain a weight update rule of

$$w_i \leftarrow w_i + \eta \sum_{d \in D} (t_d - o_d) x_{id}$$

For a multi layer network with multiple output units the back propagation algorithm needs to be used. The error needs to be summed over all the network output units

$$E(\vec{w}) = \frac{1}{2} \sum_{d \in D} \sum_{k \in \text{outputs}} (t_{kd} - \sigma_{kd})^2$$

The back-propagation algorithm to determine the weights of a feed-forward neural network is given in appendix A

3.3.2 *Composition*

Neural networks have been used in music classification, in genetic algorithms as fitness functions and more complex neural networks have even been used in algorithmic music compositions.

Simple feed-forward neural networks do not contain a mechanism to remember past history.

In [3] Mozer created a recurrent connectionist neural network **CONCERT**[2][3]. The system works as follows:

- The network is trained on sample melodies from which it learns melodic and phrase constraints
- Representations of pitch, duration and harmonic structure are included that are based on psychological studies of human perception, based on Laden and Keefe's work [28]

The system yielded good results on simple structured artificial sequences however the system performed poorly on natural music[4]. Mozer described the system as lacking musical coherency [3]. Furthermore the system performs poorly as the length of the pieces increases.

Mozer stated the reason for failure is likely due to the Recurrent Neural Network (RNN) not being able to track more distant events that build global structure [3] however a LSTM recurrent network is able to achieve this goal

---

2 CONCERT - connectionist composer of ERudite tunes
3 The ER may also be read as ERratic or ERsatz
4 One critic described the resultant melodies as compositions only a mother could love

[1].

In order to solve the problem of global structure Douglas and Jurgen attempted to use a LSTM network to compose musical pieces [1]. In this attempt the network was successfully able to learn a form a blues music and stay close to the relevant structure. The system used cross entropy as the error rate:

$$E_i = -t_i \ln(y_i) - (1 - t_i) \ln(1 - y_i)$$

where y_i is the output activation and t_i the target value for the $i-$th output unit. The topology of the network was arranged as follows:

1. Four cell blocks are connected to the input units for chords
2. The last four cell blocks are connected to the inputs units for melody
3. Chord cell blocks have recurrent connections to themselves and melody cell blocks
4. Melody cell blocks have recurrent connections to other melody cell blocks
5. Output units for chords are connected to cell blocks for chords and to input units for chords
6. Output units for melody are connected to cell blocks for melody and to input units for melody

The underlaying chord structure was kept fixed.

The results indicate that a LSTM network is able to compose with both local- and global structure from a set of training data [1].

3.3.3 *Conclusion*

Simple feed-forward neural networks do not have the functionality to remember past history and thus do not have the capability to evaluate repetitive rhythmic patterns.

Recurrent neural networks are able to encode temporal information though initial investigation by Mozer lead to limited success as compositions lacked global structure.

Further investigations by Douglas and Jurgen indicated that LSTM networks are able to compose with local and global structure.

3.4 GENETIC ALGORITHMS

3.4.1 *Background*

Genetic Programming, not to be confused with evolutionary or genetic algorithms is a evolutionary algorithm based methodology to find computer programs that perform defined tasks by simplistically mirroring biological

evolution. The more fit programs carry on their chromosomes into future populations. Fitness is rated by a fitness function. Other genetic operators such as recombination and crossover are also usually applied. These evolved genetic programs are usually represented in tree form. Figure 10 indicates the function $(2.2) - (\frac{X}{11}) + 7\cos(Y)$ written in tree form.

$$\left(2.2 - \left(\frac{X}{11}\right)\right) + \left(7 * \cos(Y)\right)$$

Figure 10: Figure of a example genetic program tree

A genetic algorithm consists of the following components:

1. Representation for chromosomes[5]
2. Initial population of chromosomes
3. A set of genetic operators to alter the population
4. A fitness function to assess the fitness of an individual
5. A selection method to determine which individuals in a population survive

The algorithm proceeds as follows:

1. Initial population is randomly generated
2. The fitness of each individual is assessed
3. Individuals are selected to which genetic operators are applied, e.g.
   Two parents are selected to generate a new child with crossover
   Random mutation occurs in individual
4. Various forms of selection are available that determine which individuals will be in the next generation

Genetic operators are used to generate diversity A genetic algorithm has a fixed set of genetic operators. Operators may include:

- Reproduction - A parent from the population is carried over to the next generation
- Crossover - genotype of both parents are combined using different procedures

---

5 Also commonly referred to as an individuals

- Mutation - A single mutation is applied to a chromosome at a set mutation rate

The choice of a fitness function is a big problem when using a genetic algorithm to compose music [29, 30]. There is no objective method to rate whether a melody is good or bad [31]. Traditionally when posed against such tasks the fitness function is provided interactively by the user, i.e. the user rates whether the piece is good or bad

The interactive GA approach is an approach to the fitness function where a human interactively rates the quality of a composition (fitness). A well known software that utilizes a neural network and a interactive interaction for a fitness function is GenJam [9]. The drawback of having the user interactively evaluate the fitness of individuals is that it is time consuming and poses a processing bottleneck [1]

Selection is the choice of which individuals will be chosen for the next generation. Selection concerns the reproduction and crossover operators. Some methods of selection include:

- Roulette wheel selection - chance of individual being chosen is proportional to fitness
- Tournament selection - tournament is staged between two individuals to determine which one gets selected[6].
- Elitist selection - commonly the best few individuals are kept unchanged but are eligible to be selected as parents.

3.4.2 *Composition*

Genetic algorithms have already been used in a variety of work in algorithmic music composition. Some of these include:

- Thematic bridging [32]
- Composition systems using IGAs [33]
- IGAs to improvise jazz solos [9, 4]
- Integration between interactive genetic algorithms and genetic programs [34]
- Hybrid approaches employing statistical, connectionist and evolutionary elements [11]
- Various work into different fitness functions

In Alfonseca's system the generated musical pieces had the style of well-known authors even when the fitness function only took relative pitch envelope into account and all generated note lengths were of fixed duration [35].

Thematic bridging is the application of an initial music pattern to a final pattern over a specified duration [32]. In this approach Horner modified

---

6 This is tournament selection in its simplest form

or reordered elements in a music pattern through various operations. For example:

Given the initial note pattern of:

`Gb Bb F Ab Db`

and a final note pattern of:

`F Ab Eb`

the musical output could be:

`Gb Bb F Ab Bb F Ab Gb Bb F Ab Eb F Ab F Ab Eb`

by means of various operations such as mutation, rotation, deletion and so on. For thematic bridging a composite fitness function was used which rates how close the developed pattern matches the final pattern and whether the ordering of the elements are correct

Similarly a system of variations was developed Jacob that proceeds as follows [21]:

1. Define a primary set of motives
2. Compose phrases by layering and sequencing new motives
3. New motives are created by variations of motives already in the phrase
4. Phrases are combined together

Jacob's system had a human judge evaluate the individual chromosomes.

NCD has also been commonly used as a fitness function, for more information about the normalized compression distance see sections 4.2 (in music classification) and 5.5 (as a fitness function). Alfonseca proposed the following genetic algorithm scheme for composing melodies [35]:

1. Define a set of M musical pieces for a guide set G
2. Encode both the guides and individuals in the population as pairs of integers where the first integer represents the note interval and the second the length as a multiple of the minimum unit of time.
3. See eq. 5 on page 34 for the fitness function used
4. The 16 lowest fitness genotypes of every generation is removed
5. The 16 highest fitness genotypes of every generation are paired by means of genetic operations.

3.4.3 *Conclusion*

There are a variety of alterations on genetic algorithms and genetic programs, however evolutionary algorithms in general are a viable means to compose music.

The encoding of a musical piece as a chromosome affects the interactions of the genetic operators on the musical piece and most authors encode the problem differently.

It is important to restrict the domain of problem otherwise the search space for the genetic algorithm may be too large [21]. Most of the studies listed in this document had restricted goals. For example, using only two octaves for the notes significantly reduces the size of the search space and many real melodies comply with it [35]

The fitness function is an important part for having the genetic algorithm result to good melodies. See section 5 for insight literature has on fitness function for evolutionary algorithms. The representation of melodies for the algorithm is arguably just as important.

For this project the focus is on evolutionary and learning algorithms and as such other procedural means of music composition will be neglected.

Fitness functions for genetic algorithms are considered in chapter 5. Possible music classification algorithms will also be investigated as they can serve as fitness functions.

MUSIC CLASSIFICATION

4.1 INTRODUCTION

In this section we will investigate methods of music classification. If some algorithm is a good method to rate the closeness of a song to a genre or style then it may also serve as a good fitness function for a evolutionary algorithm.

4.2 NORMALIZED COMPRESSION DISTANCE

4.2.1 *Background*

The Kolmogorov complexity of piece of text is the measure of the computable resources needed to specify the text. The complexity of a string is the length of the shortest possible description of the string in a fixed universal description language [36].

The information distance between two string x and y is defined as the length of the shortest program p that can compute x from y and y from x The length of p can be expressed using Kolmogorov complexity [37]:

$$|p| = \max\{K(x|y), K(y|x)\}$$

The information distance p is a absolute measure. A more useful similarity metric is one that expresses the distance in relative terms. The Normalized Information Distance (NID) is given by [38] :

$$NID(x, y) = \frac{\max\{K(x|y), K(y|x)\}}{\max\{K(x), K(y)\}}$$

The concept of NID is important, however it is not computable. An approximation of the normalized information distance is commonly used. $K(x)$ is approximated by $Z(x)$ where $Z(x)$ is the binary length of a data x compressed by a compressor Z.

$$NCD(x, y) = \frac{Z(xy) - \min\{Z(x), Z(y)\}}{\max\{Z(x), Z(y)\}}$$

where $Z(xy)$ is the length of $x + y$ compressed by Z. Any good compressor may be used for Z such as

- Gzip
- Bzip2
- Lempel-Ziv and its variations

4.2.2 *Literature*

Normalized compression distance has been used in a variety of cases. It has been used in applications of general clustering and classification of data in arbitrary domains. This includes music classification [33].

Cilibrasi and Vitiyani used the Normalized Compression Distance to approximate the Kolmogorov Distance between different musical pieces as a method to compute clusters of music [33]. The MIDI files were pre-processed such that when two notes occur at the same time only the note with the highest pitch is kept. The music was represented as a string and the distances between different musical pieces was computed.

Ctaltepe, Sonmez and Adali used the normalized compression distance and used it to classify music pieces using k-nearest neighbors [39]. The training data has a label associated. The closest k training data (by NCD) to a song is obtained and the most frequent label in the k set is used to classify the musical piece. Ctaltepe Sonmez and Adali found that the distance measure works better when more training data is available and the performance is dependent on how the input data is pre processed. The best results were obtained when the midi files were sampled at 1ms and the k = 1 nearest neighbor identification was used. The music was represented in the following format: outputting the first note and then the difference in pitch between consecutive notes. Using the above means a classification accuracy of 79% was achieved on 57 midi files.

Li and Sleep have also found that the 1-nearest neighbor with a Lempel-Zip compressor outperformed more complex statistical methods and compressors. Using relative pitch intervals in the music representation outperformed using absolute pitches. A performance of 92.35% was obtained [19]. The midi files were organized into four categories Beethoven (302 files), Haydn (261 files), Chinese (80 files) and Jazz (128 files). The dataset is unbalanced and the study does not make it clear which partition of the dataset was used as training samples and what partition was used for verification.

In [35, 40, 19, 7] it was found that the normalized compression distance serves as a promising fitness function for genetic algorithms for automatic music generation. Thus NCD may be viably used as a fitness function for a genetic algorithm and as a metric to help classify music.

4.3 NEURAL NETWORKS

McKay investigate using K-nearest neighbor techniques and artificial neural networks in order to classify MIDI music by genre. He included some of the following metrics as input to the neural network [41]:

- Number of notes - standard deviation of number of notes activated in each channel
- Note duration - standard deviation of total duration of notes
- Dynamics - standard deviation of average volume of notes

- Melodic Intervals - average melodic interval
- Simultaneity - average number of notes that are played concurrently
- Note density - average number of notes per second
- Average time between attacks - average time between note activations
- Initial tempo - tempo in beats per minute
- Pitch variety - number of pitches used at least once
- Most common pitch class - Most common pitch divided by number of possible pitches

The complete set can be found in Mckay's dissertation [41] Using the complete list of metrics neural networks had a success rate of 83%.

FITNESS FUNCTIONS

5.1 INTRODUCTION

In this section we will briefly review some functions that may be used as a fitness function for a genetic algorithm.

Algorithms that help classify music could possibly be used as fitness functions, as this would rate the similarity of the evolved piece of music to a genre or style. These algorithms are found in section 4. If work has been done that has used the music classification algorithm as a fitness function it will be listed in this section.

Fitness functions can be categorized as follows:

- Interactive - The user provides the fitness of a piece of music
- Expert systems - The musical fitness is assigned according to best practices commonly studied in music theory
- Learning based functions - Fitness functions that learn from previous data

Learned fitness functions are sensitive to input data and the selection of features is a difficult task [8], however the focus here is on learning algorithms.

5.2 ZIPF'S LAW

Zipf's law states that the frequency of an event is inversely proportional to its statistical rank, that is:

$$f \propto r^{-a}$$

where f is the frequency of occurrence of a particular event and r is the statistical rank. a is close to 1. Zipf's law can also be stated as:

$$P(f) \approx \frac{1}{f^n}$$

where $P(f)$ denotes the probability of an event of rank f and n is close to 1

One can determine whether melodic intervals follow Zipf's law by counting the melodic intervals in a piece. The result is usually plotted on a logarithmic scale and is known as a rank-frequency distribution

Zipf found evidence for the theory in music. An analysis of Mozart's Bassoon Concerto in Bb Major revealed an inverse relationship between the length of intervals between repetitions of notes and the frequency of their occurrence [42].

Several other different rank frequency distributions can be obtained for a musical piece. These include [12]:

- Pitch - distribution of MIDI pitches
- Chromatic tone - distribution of 12 chromatic tones
- Duration - note durations
- Melodic interval - distribution of melodic intervals
- Melodic bigrams - distribution of adjacent melodic interval pairs

Linear regression is performed to obtain the slope of the distribution. The coefficient of determination R^2 is also computed in order to determine how well the slope fits the data.

Figure 11 plots the rank frequency distributions of various metrics for the Beatles' song Let It Be. The figures were generated by Jensen for his thesis [8]. The metrics can be seen to follow a Zipfian distribution. Results by Manaris also indicate that most music pieces display near Zipfian distributions [12].

Zipf-based metrics capture essential parts of scaling properties in music. These metrics indicate that music follows a distribution balanced between near-zero slope and steep negative slope. Different styles of music have different slopes. There exists also a correlation between Zipf metrics and human preference [12].

Jensen used a Gaussian to define the target fitness as [8]:

$$f_m(a, b) = e^{(\frac{b-a}{-\lambda})^2}$$

where a is the metric slope of an evolved piece of music, b is the target slope and λ is the tolerance

Since there are several metrics for a given piece of music the fitness function should incorporate these. Jensen used the weighted sum of several metrics.

$$f(\vec{a}, \vec{b}) = \sum_{i=1}^{N} w_i f_i(a_i, b_i)$$

Jensen has found that Zipf metrics can be used to evolve pleasant music using a tree-based representation, however the majority of the evolved melodies were unpleasant [8]. Zipf metrics only capture the scaling properties of distributions and ignore the musical events that account for different frequencies. Zipf's law neglects musical content and can be seen as knowledge weak. Jensen concluded that the Zipf metrics are insufficient for musical fitness alone.

Manaris had more success with Zipf metrics however he used them as input to an artificial neural network to evaluate the fitness of melodies, however Manaris states it is wiser to use the fitness function in a partially interactive system [12]

---

1 Figure generated by Jensen for his thesis [8]

Figure 11: Rank frequency distributions and slopes of different metrics for The Beatles' Let It Be [1]

5.3 COSINE SIMILARITY

In Information Retrieval cosine similarity is commonly used to asses the similarity of two documents:

$$\text{sim}(\vec{A}, \vec{B}) = \cos\theta = \frac{\vec{A} \cdot \vec{B}}{|\vec{A}||\vec{B}|}$$

where $\vec{A}$ and $\vec{B}$ are two document vectors

In order to rate the similarity between music scores features such as pitches and melodic intervals are used.

As with Zipf's law the fitness is the weighted sum of similarity measures:

$$f(\vec{A_1}, \ldots, \vec{A_n}; \vec{B_1}, \ldots, \vec{B_n}) = \sum_{i=1}^{N} w_i f_i(\vec{A_i}, \vec{B_i}) \tag{4}$$

The fitness function rates the fitness of the evolved individual with a target piece.

Jensen conducted multiple experiments using the fitness function in equation 4. The metrics covered indicate the frequency of certain events. At first only a single metric was included and thereafter multiple metrics. As more metrics were included the evolved melodies became more similar to the target piece. More pleasant melodies were evolved when the target piece was The Beatle's Let It Go than Mozart's Piano Sonato No. 16. There was no correlation between melody and rhythm. Metrics included for the fitness function were:

- Pitch
- Chromatic tone
- Melodic interval
- Melodic bigram
- Rhythmic interval
- Rhythmic bigram

Jensen concluded that the results obtained by the cosine similarity fitness function were more pleasant than those obtained by Zipf's law as Zipf's law rates music on scaling properties only [8].

5.4 NEURAL NETWORKS

Different forms of networks have been used as a fitness function for evolutionary algorithms. Some of these include:

- Adaptive resonance theory neural networks using binary classification patterns [10]
- Recurrent neural networks

- Cascade correlation neural network designed to reduce GenJam bottleneck [9]

Some common problems with neural networks as fitness functions is that they require a lot of time to be trained, require a good representation for a set of inputs to map to an output and the structure of the neural network is fixed after training [10].

In [9] Biles tried to design a cascore neural network to rate musical scores. Since a neural network outputs fitness based on the input parameters the choice of input metrics are important. Metrics that included the number of new note events in a measure, the number of unique new note events, the size of the maximum interval, the number of changes in a direction between adjacent notes failed to capture the fitness for the Artificial Neural Network (ANN) [9].

Biles argues that the reason for this is that humans listen to music in more complex ways and that simple statistical measures fail to capture this. Zipfs law, which only captures the scaling properties of music also yielded poor results as a fitness function for similar reasons [8] (See section 5.2).

An Adapative Resonance Theory (ART) network has been used as a fitness function whereby a Genetic Algorithm (GA) utilizes clustered representations of rhythm styles to interactively generate rhythm patterns to according to a certain style [10].

An adaptive resonance theory network utilizes unsupervised learning and clustering algorithms to recognize patterns. New clusters are created if a pattern cannot be associated with existing clusters. Another characteristic of ART networks is that new training does not cause loss or corruption of old training data [43]

A ART1 network clusters binary vectors. The basic structure of an ART1 network involves (See figure 12) :

1. Input processing field - F1
2. Cluster units - F2
3. Mechanism to control degree of similarity of patterns in the same pattern
4. weighted bottom-up connections between F1 and F2 layers
5. weighted top-down connections between F2 and F1 layers

Burton had the ART network fitness function operate as follows [10]:

1. Each individual in the population is an input to the ANN
2. The network determines the winning cluster
3. The degree of similarity between the individual and the cluster is determined
4. If the degree of similarity is above a certain threshold the individual is added to the cluster

Figure 12: Figure of ART ANN topology

5. If no clusters match the individual closely enough a new cluster is created

6. Fitness is assigned as a degree of similarity to a cluster.

NEUROGEN is another attempt at using a neural network as a fitness function to compose small diatonic, western, four part harmony compositions. The system has shown limited success however it was able to produce 4 bars of music [2].

Chen used a Simple Recurrent Network (SRN) with composition rules on tonality and rhythm as a fitness function for a GA [6]. The simple recurrent network has an input layer that represents a measure at time T with the output layer representing the measure at time $T + 1$. A recurrent network is used as a single step predictor to compose music. The network predicts notes at $T + 1$ using the notes at time T. After the network has been trained it can be seeded with inital values to generate novel compositions. The following constraints were used:

1. Pitch diversity constraint - number of measures with unique pitch sequences

2. Rhythmic diversity constraint - number of measures with unique signatures

3. Measure density constraint - ratio of number of notes to maximum notes

4. Pentatonic pitch class constraint - number of notes that belong to pentatonic pitch class

5. Cell structure - ratio number of times cell pattern occurs to maximum number of patterns

The system was able to generate melodies with systematic structure however it lacks global structure. Individual measures sound pleasant and diverse however there is a lacking structure as a whole

5.5 NORMALIZED COMPRESSION DISTANCE

As noted in section 4 the Normalized Compression Distance has been used to help classify music genres. However Normalized Compression Distance has been explored as a possible fitness function for evolutionary algorithms [35, 40, 7]

The fitness function used by Alfonseca et. al for an individual x and a guide set G was defined as [7]:

$$
f(x) = \left(\sum_{g_i \in G} \mathrm{NCD}(x, g_i) \right)^{-1}
\tag{5}
$$

Where g_i is a guide in guide set G containing M musical pieces and x is the set of differences between consecutive notes.

Alfonseca encoded the chromosomes as N vectors containing a pair of integers. The first integer denotes the note interval and the second represents the length as a multiple of the smallest unit of time [35].

5.6 INTERACTIVE FITNESS FUNCTIONS

Genetic algorithms which employ user interaction as a means of rating the fitness of are known as Interactive Genetic Algorithms (IGAs).

Unehara constructed a system that composes 16-bars music using a GA [13]. The user rates individual chromosomes, new chromosomes are applied by genetic algorithm and the user is asked to rate the individuals again. Should the user find a good piece they may favorite it. The fitness of the chromosomes were seen to increase as the generations increased, however it is unknown whether the melodies were pleasant.

Using an interactive fitness function may lead to better results than most other fitness functions however it is a tedious and demanding process and may also lead to inconsistencies in evaluation. Some researches try to reduce this effect by constructing ANNs which learn the user's ratings, and as such may be used in place of the interactive fitness function [9, 14]. See section 5.4 for the use of ANNs as fitness functions.

Johnson and Poli had a user rate individual sequences and trained a neural network base automatic rater, which may replace the user in larger runs. The musical pieces generated by the automatic rater were pleasant but they were not as pleasant as the musical pieces generated by the user interactive runs [15]

The superiority of a interactive fitness can be seen, as a person can rate the pleasantness, or fitness of a song much more accurately than current

quantitative fitness functions. However this imposes a bottleneck on the system as it is time consuming. A partially interactive system may yield a good compromise [9].

CONCLUSION

We have investigated numerous methods to compose music algorithmically. The two most prominent methods currently to compose music is through genetic algorithms and neural networks.

Genetic algorithms require proper music representation and a good fitness function. Multiple work has gone into investigating various possible fitness functions. The three most promising candidates are NCD, ANN and cosine similarity.

Interactive genetic algorithms yield good results although this imposes a performance bottleneck on the system, as the system is required to wait for user input. This is a time-consuming process although a partial interactive system might be viable.

Simple feed-forward neural networks are unable to compose music due to their inability to encode temporal information. Recurrent neural networks were investigated and early findings yielded poor results, though LTSM networks seem promising.

Algorithmic music composition is viable although there is large room for improvements to be made. Currently only short musical pieces sound pleasant. Longer pieces tend to be repetitive and lack global structure.

Most current methods restrict their domain in order to investigate only the main research questions.

A hybrid approach may yield good results.

Part III

CONCEPTUAL DESIGN

INTRODUCTION

In this part a conceptual design is proposed to solve the problem of composing music algorithmically. Composing music algorithmically can be done in a variety of ways. Different types of algorithms can be utilized:

- Algorithms that utilize expert knowledge. Expert systems utilize knowledge from a specific domain in order to make decisions.
- Algorithms that learn. These algorithms utilize pattern recognition to learn from data in order to make decisions.

For an expert system to be used in algorithmic music composition knowledge from music theory is utilized in order for the algorithm to make decisions. When too many constraints are imposed on the algorithm the variety of the type of music produced is lessened.

For this project however, the focus is on algorithms that can learn. This direction is taken as the application is required to produce music in various different styles. In order to design an application that utilizes machine learning algorithms to compose music the project will first be decomposed into discrete units. The function and interaction of each unit will be investigated.

Different types of machine learning algorithms will be investigated and the best solution will be chosen that satisfies the requirements of the project.

SYSTEM ARCHITECTURE

In this section the project will be broken down into discrete functional components and the interaction between different components will be investigated.

Figure 14 shows the interaction of the user with the application. The user interacts with the application and selects a certain style of music to be composed. The application generates the music according to a certain algorithm and plays back the generated piece.

A user is able to select a certain style of music for composition. Once the piece is generated the user is able to save the music to a MIDI file.

The core functionality of the application resides in the algorithm that is responsible for composing a piece of music.

By incorporating all these elements we can concretely describe the project in discrete components in a functional architecture diagram, see figure 13.

8.1 FUNCTIONAL ARCHITECTURE

8.1.1 *Functional Unit 1 - Operator*

The operator is responsible for interacting with the application. The operator will select the style or category for composition and instruct to application to compose music. If a certain piece of music is to be save the operator will be
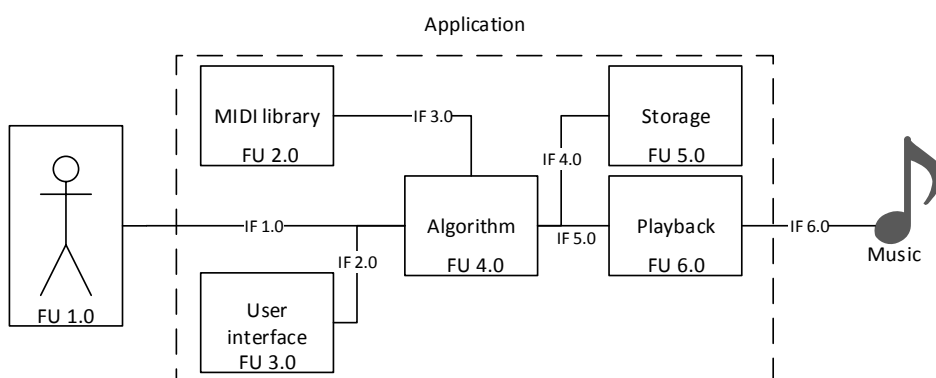
Figure 13: Functional architecture

responsible for instructing the application to do so and to select the location the MIDI file is to be saved

8.1.2 *Functional Unit 2 - MIDI library*

The MIDI library is a large collection of MIDI files. The files in the MIDI library are organized into categories. Each category represents a certain style of music. The algorithm will utilize a subset of the library (a category) as an input dataset.

8.1.3 *Functional Unit 3 - User Interface*

The user interface is the front end of the application. The operator interacts with this functional unit in order to instruct the application what to do.

The user interface should be designed in a user-friendly manner in order to accommodate the operator.

This user interface allows the user to select a certain style, instruct the algorithm to compose music and to save or play back a piece of music once it is composed.

8.1.4 *Functional Unit 4 - Algorithm*

The algorithm is the central part of this project. The algorithm converts the input MIDIs into output MIDI.

The algorithm will utilize a category of MIDI files from the MIDI library in order to compose a new piece of music not in the MIDI library.

The output piece of music will represent the style of music that was used as input.

Various machine learning algorithms are investigated. The results of the investigation is covered in part iv.

8.1.5 *Functional Unit 5 - Storage*

This unit is responsible for storing the output MIDI from the algorithm into a MIDI file.

8.1.6 *Functional Unit 6 - Playback*

This unit is responsible for playing back the output MIDI from the algorithm through an audio output device.

8.2 INTERFACES

The interfaces indicate the interaction between different functional units. The interfaces have the following functions:

1. Interface between the user and the user interface. This input would be from a input peripheral device.

2. Interface between the user interface and the algorithm. The interface calls the algorithm with the parameters supplied by the user such as when to start and what type of style was selected

3. Interface between the MIDI library and the algorithm. The input to the algorithm is a selection of MIDIs. The interface converts MIDI files into a format required by the algorithm

4. Interface between the algorithm and storage. This interface converts the output from the algorithm into the MIDI file format

5. Interface between the algorithm and playback. This interface converts the output from the algorithm into a format that is ready for playback through the playback functional unit.

6. Interface between the playback and the music. This represents the audio output device and it's workings.

OPERATIONAL FLOW

The operational flow indicates the interaction between the operator and the application and how the operator should use it.
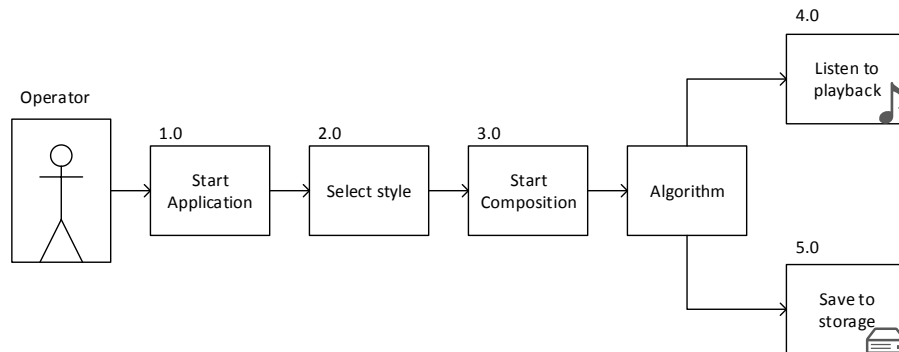
Figure 14: Operational Flow

Figure 14 shows how the operator interacts with the application.

For this application, the operator selects the style of music to be composed; instructs the application to compose the music according to the selected style and then to instruct the application to either play back the piece of generated music or save it to storage.

Figure 15 shows a more in-depth interaction between the user and the application. The application allows the user to:

1. Select the style for composition
2. Select the algorithm to be used for composition
3. Load a previously generated melody
4. Save a generated melody
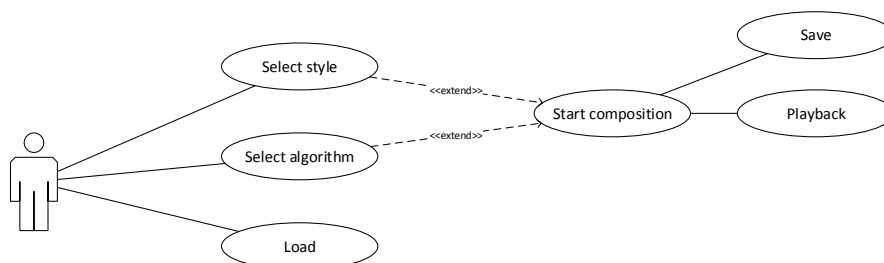
Figure 15: Application interaction use case

Some considerations need to be taken to ensure that the interface to the application is as user-friendly as possible.

1. The design of the interface should be clean and minimalistic.
2. Clear instructions should be given to the user when there is a possibility of ambiguity.
3. Images and icons should be used as visual aids for buttons or to indicate the function of a feature
4. Separate functional features should be separated in the interface. The following areas should have seperate section in the interface
   - Playback, saving and loading
   - Generation of melodies
   - Optional options and parameters for algorithms

Functionally, it makes sense for the user to choose the filename and path of the melody to be saved. Similarly the user should be able to select which previously generated melody should be loaded.

CONCLUSION

In order to solve the problem of generating music algorithmically the task was broken down into its functional architecture. From this each functional unit and its interaction is made apparent. A clear separation was made between the different components and work can be done on each separately.

The core component of this project is the algorithm which is used to compose music. Possible machine learning algorithms will be investigated and implemented. Algorithms yielding good results will be kept for inclusion into the application. Possible music categories which serve as input datasets to the algorithms will also be investigated.

Part IV

ALGORITHMS

INTRODUCTION

In this part possible algorithms will be investigated, designed for composition and implemented in order to determine which are viable for inclusion into the application.

The system operates on MIDI files. MIDI events are processed and the resulting notes are stored in a sequence where the pitch and duration of each note are recorded. A note is represented in MIDI as a number from 0 to 127. Thus for each MIDI track a monophonic melody can be extracted.

A large set of MIDI files (600) were collected for style of music. The following categories were used to test the algorithms:

1. A mixture of Classical music (665)
2. A mixture of retro video game music (395)
3. A mixture of Pop and Top40 (1370)
4. A mixture of Classic Rock (2274)
5. A mixture of Jazz music (114)
6. A mixture of Dance and Techno (214)

These categories are not the final list of styles that will be included into the end user application.

REPRESENTATION

This section discusses how music is represented internally in the system and how the data structures are designed. Proper data structures are extremely important for good algorithms [1].

12.1 ASSUMPTIONS AND SIMPLIFICATIONS

For this system our model of music is simplified and we do not make provisions for chords. Melodies are monophonic and are simply a list of notes of their durations.

Furthermore, we assume all files are MIDI *Track type* 1.

12.2 NOTE

A note is the most basic element. A note consists of:

1. Note duration - A whole note, half note, quarter note and so on.
2. Note pitch - A number representing the note pitch, in MIDI scale. $0 \leqslant P < 128, p \in \mathbb{Z}$

From these properties the following can be derived:

1. Octave - Every twelve note pitches represent one octave
2. Chromatic tone - An integer from 0 to 11. Also corresponds to the name of the note for example C, C#, G, A#

These properties are related to the Note pitch by:

$$P = 12o + p$$

where o is the octave and p is the chromatic tone

12.3 MELODY SEQUENCE

A melody sequence is simply a list or sequence of notes:

$$\{N_1, N_2, \ldots, N_n\}$$

Where N is a note.

Most algorithms operate only on the melody sequence and output a melody sequence. Some algorithms may optionally require information on the instrument in which the melody sequence is played.

---

1 Bad programmers worry about the code. Good programmers worry about data structures and their relationships ~Linus Torvalds

Figure 16: Illustration of a composition and its elements
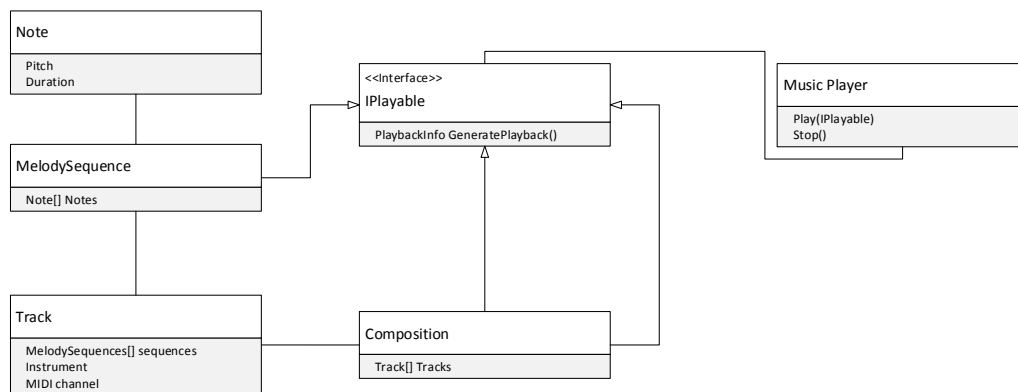
Figure 17: UML diagram indicating the interaction between compositional elements

12.4 TRACK

A track contains a sequence of melody sequences and is also described by an instrument in which the sequence are played.

12.5 COMPOSITION

A composition is a sequence of tracks which can be played in parallel. MIDI files are converted into compositions.

Figure 16 illustrates the relation between the above mentioned structures for a midi file.

Figure 17 indicates the interactions between Notes, Melody Sequences, Tracks, Compositions and the Music Player.
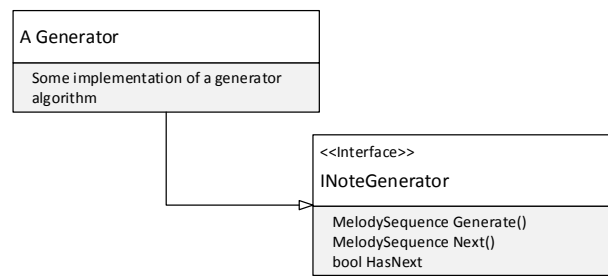
Figure 18: UML diagram of a generator class

12.6 GENERATOR ALGORITHMS

All algorithms considered in this section derive from the interface INote-Generator. This exposes basic functionality for generating a algorithm and determining whether the Next function can be called. The Next function returns a new random melody sequence without resetting the generator. See figure 18.

Using this abstracted design allows the user to select any algorithm to be used for composition, and the main application logic only needs to call Generate() or Next() as required. The details of the algorithm implementation is abstracted away.

GENETIC ALGORITHM FOR MELODY GENERATION

For this project a quantitative approach is taken toward algorithmic music composition. In particular quantitative metrics will be used in the fitness functions of the genetic algorithm.

In section the following types of music composition algorithms were investigated:

1. Neural Networks
2. Genetic Algorithms

In the literature review, it was found that the pieces generated by neural networks lack musical coherency and perform poorly as the length of music increases. Some other attempts have met slightly more success although the overarching view for neural networks in music composition seems grim[1].

Genetic algorithms are viable for algorithmic composition for the following reasons:

1. They allow for great flexibility in implementation and music representation
2. There has been some research into possible fitness functions for evolutionary music composition.
3. Great amount of variety made possible by different fitness functions and by the representation of music used in GAs

In section 5 the following fitness functions were covered:

1. Zipf's law
2. Cosine similarity
3. Neural Networks
4. Normalized Compression Distance
5. Interactive evaluation[2]

Figure 19 indicates the flow for a genetic algorithm/program.

13.1 INTRODUCTION

Evolutionary algorithms come in a variety of different types. The two most common types found are Genetic Algorithms and Genetic Programming. Genetic Programming has been found to be more suitable for composing music than genetic programming due to music forming a hierarchical structure.

---

1 Although neural networks as functions in genetic algorithms have had better success
2 An interactive fitness function imposes a bottleneck on the performance of the system
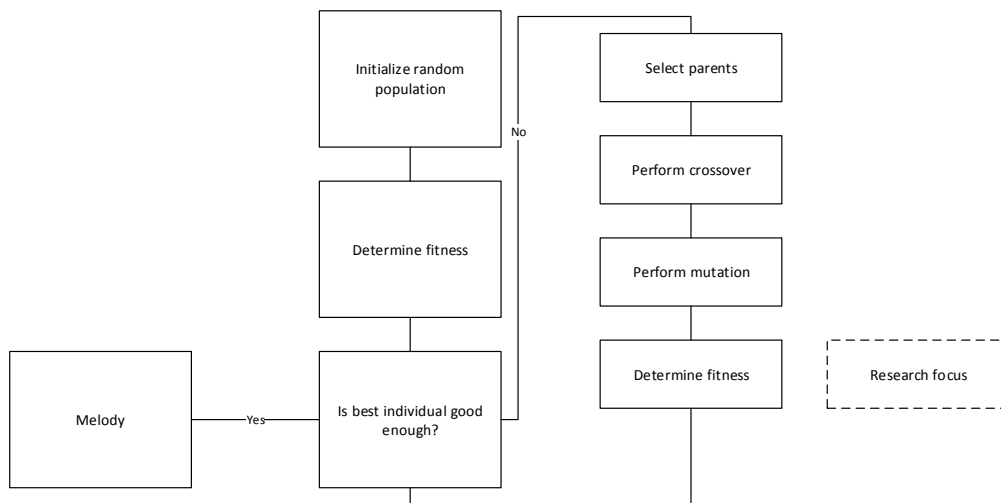
Figure 19: Flow diagram of the operation of a genetic algorithm

A flexible genetic programming model will be developed that is able to function with the investigated fitness functions.

The two largest problems for a evolutionary computing problem is:

1. Obtaining a good representation of the problem
2. Obtaining a good fitness function

A genetic algorithm modifies the structure of an individual, if the structure is poorly chosen then a optimal solution wont by found by the algorithm. An ideal fitness function is able to quantify the fitness of an individual. An ideal fitness function in music composition would map the human perception of pleasantness into a fitness value.

13.2 FITNESS FUNCTIONS

The fitness function is a primary research interest in genetic music composition. An ideal fitness function captures the human perception of pleasantness in music, however there have not been practical advances in this direction. Statistical measures will rather be used.

Of the set of investigated fitness functions only the following functions will be implemented:

1. Cosine similarity
2. Normalized Compressions Distance

Since interactive evaluation is slow, Zipf's law is superseded by Cosine similarity and training neural networks as fitness functions pose a large number of problems.
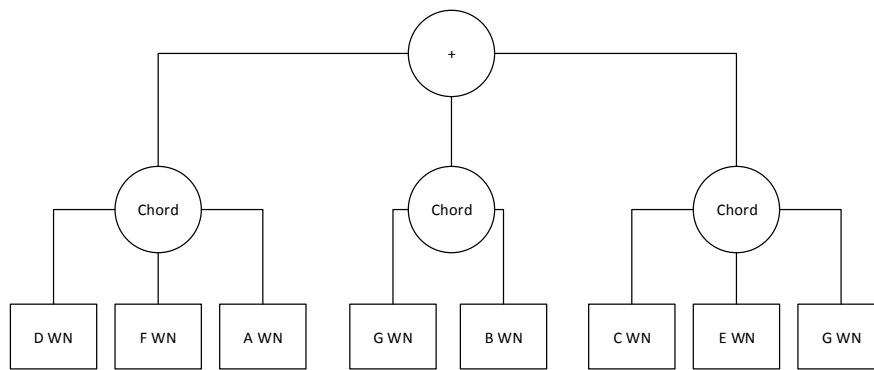
Figure 20: A music piece in a tree structure

13.3 MUSIC REPRESENTATION

A flexible music representation can also be developed for the Genetic Programming algorithm.

Figure 20 indicates a polyphonic melody in tree form. Strictly a chord is a set of three or more notes that are played simultaneously. Note that the second branch is only constituted of two notes.

Minsky and Laske [23] argued for a tree representation of music since the tree represents the hierarchical nature of music. The tree representation is much more complex than data structures such as vectors that are used in GAs.

Some authors [4] limit the search space by ensuring that melodies are in a certain scale.

The representation will model MIDI events and manipulation of them. A simpler tree structure will be developed as chords will not be supported by the system. The operators of the tree will be given in the next section.

13.4 GENETIC OPERATORS

Several different operators can be performed on the tree structure. In figure 20 the serial concatenation and parallel concatenation operators were shown. Some additional operators which may be employed include:

- Repetition - Repeating a segment a given number of times
- Shift note pitches - Shift all pitches of notes by a certain amount
- Duration elongation or contraction - For example a slow operation doubles the duration
- Transposition - Moving note positions relatively
- Retrograde - Reversed order of notes

Genetic operations such as mutation and crossover may be used conventionally.

13.5 METRICS

Feature extraction is required to reduce the search space and to provide the fitness function with musically meaningful measures with which to rate the fitness. In this section we cover some quantitative metrics that may be used as measures to identify or represent music.

The frequencies of a metric is used in the Cosine Similarity fitness function.

The different types of metrics that will be used include:

1. Pitch - List of pitches
2. Pitch differences - Store first pitch, thereafter list consecutive pitch differences
3. Chromatic tone - 12 pitch class. Notes are reorganized into 12 classes.
4. Note durations - durations of notes
5. Pitch distance - intervals between repetition of pitches
6. Chromatic tone distance - intervals between repetition of chromatic tones
7. Melodic interval - intervals between the current note and the previous note
8. Melodic bi-gram - Pairs of melodic intervals
9. Rhythm - duration of a note in addition to the following rest
10. Rhythmic interval - relationship between adjacent note rhythms
11. Rhythmic bi-gram - Pairs of adjacent rhythmic intervals.
12. Chromatic tone duration - A pair of the chromatic tone and the duration

Let p denote the pitch of the note, where $0 \leqslant p < 128$. Then the chromatic tone is given by:

$$c = p\%12$$

where % indicates the modulo operation. The melodic interval is given by:

$$mi_k = p_k - p_{k-1}$$

The melodic bi-gram is given by:

$$b_k = (mi_k, mi_{k+1})$$

Let r indicate the note duration with rests Then the rhythmic interval is given by:

$$ri_k = \frac{r_k}{r_{k-1}}$$

and the rhythmic bi-gram

$$rk_k = (ri_k, ri_r)$$

In this manner we can build a metric vector, let $m_i(A)$ denote a metric's value at position i for musical piece A. Then the metric vector is given by $\vec{M}_A = \{m_0(A), m_1(A), \dots, m_n(A)\}$

13.6 FITNESS FUNCTIONS

In this section two fitness functions are described which can be used for the Genetic Algorithm. NCD and Cosine Similarity are described here.

The fitness functions described in this section require the phenotype of an individual. Thus it is necessary to parse the genotype, the genetic tree of the individual down into a sequence of notes.

13.6.1 *Cosine similarity*

Some fitness functions such as Cosine similarity and Zipf's law operate on the features of music.

Cosine similarity can be applied to the metrics listed in section 24.5. Let $\vec{a_m}$ denote the vector of metric values according to a metric m for a music piece A and $\vec{b_m}$ denote the vector by m for a music piece B then the similarity between A and B is given by:

$$\text{similarity}_m(A, B) = \frac{\vec{a_m} \cdot \vec{b_m}}{|\vec{a_m}||\vec{b_m}|}$$

A set of metrics may be used. The fitness function is then given by a weighted average:

$$f = \frac{1}{N} \sum_k^N w_k \times \text{similarity}_{mk}(A, B)$$

where w is a weight assigned to metric mk.

13.6.2 *Normalized compression distance*

In order to utilize NCD both pieces being tested need to be encoded in the same way. Musical pieces may be encoded as metric vectors as listed in section 24.5. More complex metrics may be utilized however there has been no thorough investigation into this.

The NCD as an estimate to the NID was covered in section 4.2.

In order to utilize the NCD as a fitness function the following steps are taken:

1. Encode a set from the MIDI library according to a metric. Let $\Omega = \{\vec{M}_0, \vec{M}_1, \dots, \vec{M}_n\}$ for musical pieces 0 to n in the MIDI library that accord to a certain style.

2. Encode the population individual x according to the metric (Given by $\vec{M}_x$).

3. Employ the fitness function

The fitness function that will be used is:

$$f(x) = \left(\sum_{\vec{T} \in \Omega} \text{NCD}(\vec{M}_x, \vec{T}) \right)^{-1}$$

The following representation scheme was designed:

- A note is surrounded in parentheses.
- The note pitch is given by the chromatic tone symbol followed the the octave number.
- The duration is given after a hyphen. A standard note duration is used where possible otherwise the numeric value is given.
- A sequence of notes can be given by:

```
(C4-wn)(C3-qn)(C5-qn)
```

- When tracks are played the symbol || is used to denote notes playing in parallel.
- A instrument can optionally be specified at the start.

For example:

```
(Trombone (C6-3)(C5-tn)(G5-tn)(C6-3)(E5-tn)(G5-tn)(G6-3)(C5-tn)(G5-tn)(
    G6-en)(E5-tn)(G5-tn)(A6-3)(F4-tn)(F5-tn)(C5-tn)(A6-en)(F5-tn)(G4-tn)
    (G6-qn)...
```

13.7 METHODOLOGY

In this section the setup and parameters of the experiment to produce melodies with Genetic Algorithms are outlined.

In order to compose a melody with a genetic algorithm a fitness function is required to rate the fitness of an individual. The following fitness functions were investigated:

1. NCD
2. Cosine similarity

A proper representation is required. A melody is represented in tree form. The tree is given a maximum depth of $(\log_2 x) + 3$ where x is the average number of notes in a melody.

The following genetic operators were used:

1. Concatenation - The addition of two nodes
2. Duration Shift - Scale the duration with a factor

3. Pitch Shift - Move all note pitches a certain number up or down

4. Repeat - Repeat a note a certain number of times

5. Swap - Swap the positions of two notes

The following metrics were investigated: The different types of metrics that will be used include:

1. Pitch differences

2. Chromatic tone duration

3. Chromatic tone distance

4. Pitch distance

5. Melodic interval

6. Melodic bi-gram

7. Rhythmic interval

8. Rhythmic bi-gram

The duration and pitch a note may take on is limited to the available note pitches and durations in the dataset.

The parameters of the Genetic Algorithm were set as follows:

- Chromosome representation - Tree form
- Mutation rate - 0.1
- Crossover rate - 0.9
- Population - 30
- Selection - Elite selection
- Auto shuffling of individuals in population after every epoch
- Standard amount of generations - 1000

The individual with the highest fitness was used as the resulting chromosome. Since the result is in tree form it needs to be parsed into a simple sequence of notes.

13.8 RESULTS

A simple test was done comparing the NCD and Cosine Similarity fitness functions against each other. The Cosine Similarity was tested using the Chromatic Time Duration, Rhythmic Bigram and Melodic Bigram frequency metrics. The average fitness and maximum fitness were compared against each other. The genetic algorithm was run for 100 epochs.

Table 2 shows the results. It can be seen that the maximum fitness achieved by the NCD fitness function in 100 epochs is only 0.0016 whereas the Metric-Similarity is at 0.683. The pleasantness of the produced melody by NCD is also worse, however this may be due to the lower fitness. However generation time is important for the Genetic Algorithm and the NCD fitness function is slow.

Table 2: Comparison of Fitness Functions

	Performance	
	Average Fitness	Maximum Fitness
NCD	0.0016	0.00156
Metric Similarity	0.666	0.683

Table 3: Performance comparison between frequency metrics

	MaxFit	AvgFit	Time(ms)	Pleasantness
Chromatic Tone	0.944	0.9439	8992	4
Chromatic Tone Distance	0.98	0.98	11017	5
Chromatic Tone Duration	0.936	0.935	13688	5
Melodic Bigram	0.857	0.857	3900	5
Melodic Interval	0.985	0.985	9037	5
Pitch	0.832	0.832	14001	4
Pitch Distance	-	-	-	-
Rhythm	0.952	0.952	367	4
Rhythmic Bigram	0.995	0.995	4729	7
Rhythmic Interval	0.996	0.996	5501	6

Figure 21: Melody produced with a Genetic Algorithm using the NCD fitness function

Figure 22: Melody produced with a Genetic Algorithm using the cosine similarity fitness function with chromatic tone duration and rhythmic bigram metrics

The maximum fitness, average fitness, time (given in milliseconds) and subjective pleasantness were evaluated for each frequency metric. See table 3.

In order to construct a target piece for the metric similarity the notes of the first track of each midi file were sequenced. The total notes were 662923. Note that some metrics have larger time than others, this is due to the amount of features or categories for that metric and does not necessarily indicate that metric is slow. The pitch distance metric was excluded as it took too long.

From table 3 it can be seen that subjectively rhythm is important.

Figure 21 indicates an excerpt of a melody produced with the Genetic Algorithm using the NCD fitness function. The NCD fitness function produced the worst results. The resulting melodies sound chaotic and random. NCD is also slow, as each individual needs to be converted to its text representation and then be compressed in order to obtain its compressed distance.

The melody produced in figure 22 using the cosine similarity fitness function was rather pleasant and the rhythmic bigram and chromatic tone duration -frequency metrics produce pleasant melodies.

The melody produced in figure 23 used the cosine similarity fitness function using all the frequency metrics. At times the melody is pleasant however off-sounding or out of key notes are a common occurrence. The melodies also don't have structure. This indicates that some of the frequency metrics

Figure 23: Melody produced with a Genetic Algorithm using the cosine similarity fitness function with all frequency metrics

either don't work well in conjunction with each other or some frequency metrics have a negative effect on the pleasantness of melodies.

The following frequency metrics produced good or pleasant results:

1. Chromatic tone distance - intervals between repetition of chromatic tones
2. Melodic interval - intervals between the current note and the previous note
3. Rhythmic bigram - Pairs of adjacent rhythmic intervals
4. Chromatic tone duration - A pair of the chromatic tone and the duration

13.9 CONCLUSION

The decision was made to utilize a genetic programming algorithm since the tree structure accommodates the hierarchical nature of music. Genetic programming provides flexibility, variety of possible styles and a large amount of research has been done one fitness functions for genetic algorithms.

Fitness functions require good measures that make it possible to rate musical pieces quantitatively. A set of metrics were developed in which musical pieces can be measured.

Two fitness functions, namely Cosine similarity and NCD were developed to incorporate these metrics.

The results of different fitness functions were compared. NCD performed the worst. Not all of the frequency metrics required for the cosine similarity fitness function produced good results. A subset of good frequency metrics were explored.

MELODY GENERATION WITH A HIDDEN MARKOV MODEL

An attempt was made to use a hidden Markov model to generate melody sequences.

A HMM is a statistical Markov model where the assumption is made that the system being modeled is a Markov process with hidden states. See section 3.1 for more information on Hidden Markov Models for melody generation.

14.2 METHODOLOGY

In order to train the HMM, that is to re-estimate the HMM parameters the unsupervised algorithm of Baum-Welch was used. The update rules are given in section 3.1. However the regular update rules of Baum-Welch only cover an update for a single observation.

In order to operate on multiple observations the independence assumption was made, that is:

$$P(O|\lambda) = \prod_{k=1}^{K} P(O^{(k)}|\lambda)$$

For which the update rules according to Levinson are:

1. State transition:

$$a'_{mn} = \frac{\sum_{k=1}^{K} \sum_{t=1}^{T_k-1} \xi_t^{(k)}(m,n)}{\sum_{k=1}^{K} \sum_{t=1}^{T_k-1} \gamma_t^{(k)}(m)}$$

2. Observation emission:

$$b'_n(m) = \frac{\sum_{k=1}^{K} \sum_{t=1 \, mo_t^{(k)}=v_m}^{T_k-1} \gamma_t^{(k)}(n)}{\sum_{k=1}^{K} \sum_{t=1}^{T_k-1} \gamma_t^{(k)}(n)}$$

3. Initial state:

$$\pi'_n = \frac{1}{K} \sum_{k=1}^{K} \gamma_1^{(k)}(n)$$

Special care was taken to to avoid underflows. All probabilities were stored as log probabilities.

For this experiment a simple first order HMM is constructed with a ergodic structure. In addition the following constraints were made:

- The amount of states of the HMM was constrained to 100
- The duration of all notes were restricted to be in {tn, sn, en, qn, hn, wn } where wn is a whole note, qn is a quarter note and so on.

All uniques notes were linked with a unique integer id with the use of a dictionary.

In order to test the model it was trained on datasets of different styles. No distinction was made between different instruments as was the case with the Markov Chain experiment. The melody sequences of each composition was used as the input data for the modified Baum-Welch training algorithm. The number of iterations was set to 100.

14.3 RESULTS

Since the model is of first order it produces melodies without overall structure. This is to be expected.

Figure 24: Melody generated with a HMM

Figure 24 shows the melody generated with the HMM. The resulting melody sounds chaotic and lacks overall structure.

14.4 DISCUSSION

The computational complexity of the forward-backward procedure is $O(TN^2)$. Thus the training time increases greatly as the number of symbols increases and thus it is necessary to constrain the number of symbols as well as the number of states.

Some melodies, especially when the training algorithm was run for a large number of iterations were pleasant. However the results are lacking when compared with those of a higher order Markov chain. A higher order Markov chain has faster training time and the results were as good or better and no good argument could be made to choose the HMM for simple melodic generation over a chain.

The power of a HMM comes with its hidden states and observations; more complex structures benefit more from this. A better use of the model would be for harmonization or accompaniment. Allan used a HMM for harmonization of chorales [18]. The visible states represent melody notes and the hidden states chords. Such deeper structures are much more meaningful, however since the designed system does not support chords an attempt at harmonization was not made.

INSTRUMENTAL GENERATION WITH MARKOV CHAINS

15.1 INTRODUCTION

For this experiment a compositional engine will be designed that outputs melodies for a specific instrument.

In order to construct such an instrumental generator an algorithm is required that can predict time-series data.

A large advantage of recurrent neural networks over Markov chains and HMM is that neural networks have greater representational power and can take into account syntactic and semantic features. A RNN does not make the Markov assumption and is able to take into account long term dependencies. Markov chains have the advantage of being much simpler to implement and are extremely fast and efficient with the proper implementation. Recurrent Neural Networks are more difficult to implement and to train. Various training algorithms exist to train a RNN. Training a recurrent neural network is slower than a regular feed forward network. The poor results obtained from the LSTM network (See chapter 16) dissuade us from using them for the instrumental generator.

Markov Chains can also be seen as a special case of a HMM. The Markov Chain will be used over a HMM as it is simpler and the transition probabilities can be directly calculated.

This concept will utilize Markov Chains to construct a probabilistic model for a sequence of notes. The probability of the next note occurring depends on the previous state of the chain. See section 3.2.

For a Markov Chain of order m the following holds:

$$\Pr(X_n = x_n \mid X_{n-1} = x_{n-1}, X_{n-2} = x_{n-2}, \ldots, X_1 = x_1) = \qquad (6)$$

$$\Pr(X_n = x_n \mid X_{n-1} = x_{n-1}, X_{n-2} = x_{n-2}, \ldots, X_{n-m} = x_{n-m}) \qquad (7)$$

For $n > m$

15.2 METHODOLOGY

15.2.1 *General*

For this technique a different Markov Chain will be constructed for each instrument. For a certain instrument, all the notes for that instrument in a MIDI file were added to the Markov chain. This chain is used to generate notes according to a certain instrument.

A third order Markov Chain was used. That is, the previous three notes constitute the state of the model. A lower order would result in more novel

---

Algorithmus 1 : Markov Chain for a particular instrument

Data : MIDI files that accord to a certain style
Result : Markov Chain for a particular instrument
initialization;
Choose specific instrument;
for *each MIDI file* **do**

    for *each Track* **do**

        if *Track instrument is selected instrument* **then**

            add notes to Markov chain;

        ;

---

and random melodies where a higher order would produce melodies with higher similarity.

For a Markov chain the frequencies of the events can be used to calculate the transition probabilities. For a third order Markov chain the transition probabilities can be calculated by:

$$a_{i,j} = P(n_i | n_{i-3}, n_{i-2}, n_{i-1}) = \frac{\text{Count}(n_i, \{n_{i-3}, n_{i-2}, n_{i-1}\})}{\text{Count}(n_i)}$$

where $\text{Count}(n_i, \{n_{i-1}, n_{i-2}, n_{i-3}\})$ is the number of times note n_i occurred after notes $n_{i-3}, n_{i-2}, n_{i-1}$

15.2.2 *Drums*

The use of a Markov Chain can be expanded to include generation of percussion instruments.

Percussion instruments in the MIDI standard are a special case and are played on Track channel 10. Notes played on channel 10 always produce percussion sounds. Each pitch corresponds to a different percussive instrument. Table 4 displays some of these instruments. The velocity dictates how hard a percussion instrument is hit.

Table 4: Excerpt of some percussion instrument notes in the MIDI standard

Pitch	Instrument
35	Bass Drum
38	Snare Drum
41	Low Tom
45	Mid Tom
53	Ride Bell
54	Tambourine

Figure 25: Instrumental melody generated with a Markov Chain

In order to train a Markov Chain for a percussive instrument only a small modification is required to algorithm 1. In place of looping through all tracks only tracks with a channel number of 10 are used.

15.3 RESULTS

Figure 25 shows a melody generated with a Markov Chain for the Piano. A third order Markov chain was used with this melody and overall it sounds pleasant. Melodies produced with Markov Chains dont have an overarching theme but do have a local structure and theme.

Melodies produced with Markov Chains are in the same style as the input songs. This reduces compositional value and novelty, especially for higher order Markov Chains. A low order Markov Chain produces more novel although more chaotic and random melodies.

15.4 POTENTIAL IMPROVEMENTS

- Define which instruments sound good with other instruments
- Pick which instruments can be slotted together
- Ensure different tracks accompany each other and are in harmony

ACCOMPANIMENT GENERATION WITH A LSTM NETWORK

16.1 INTRODUCTION

The LSTM network is a Recurrent Neural Network (RNN), however it is more optimized and better suited to classify, process and predict time series in the case when there are long lags of indeterminate length between important events. In this experiment a LSTM network is used to generate an accompaniment track for a main melody track

A RNN has greater representational power and can take into account syntactic and semantic features. A RNN also does not make the Markov assumption and is able to take into account long term dependencies.

16.2 METHODOLOGY

The LSTM accompaniment is also generated for a specific instrument. Since it is difficult to determine which track is the main melody of a composition an assumption is made that the main melody track is the first track found within a MIDI file. The accompaniment track is another track within the MIDI file.

The LSTM network parameters were set as follows:

- Activation functions were tanh
- 2 Input units
- 15 hidden units
- 150 output units
- Learning rate of 0.1
- Momentum of 0.9

A list of all possible notes were obtained and the problem was treated as a classification problem. For a set of n possible notes the network has n possible active output units and the active output unit indicates the index of which note was activated. This was done in order to reduce the possibility of incorrect notes as even small error might cause problematic results in alternative representations (if the pitch and duration were used as output).

The input for the neural network is the note pitch and duration of each note in the melody track and the output is the index of the note activated in the accompanied track at the same index.

90% of the dataset was used for training and the last 10% was used for verification.

In order to reduce the dimensionality of the output vector filtering is applied. There are bound to be events that occur rarely which can be seen as noise, and the emphasis should be on more prominent notes. We apply a simple filtering to discard notes which have a normalized frequency below some threshold in the dataset.

$$\frac{f}{N} < k$$

where f is the frequency of the note, N the total number of notes and k the threshold. Filtering was not applied on rests.

In the first attempt rests were accounted for. In tracks with a large amount of rests, the regular notes may be filtered out.

Algorithmus 2 : Training set for LSTM network

Data : MIDI files that accord to a certain style
Result : LSTM training for a particular instrument
initialization;
for *each MIDI file* **do**
    for *each track other than melody track* **do**
        if *instrument is selected instrument* **then**
            for *each note in melody track and respective note in accompaniment track* **do**
                add normalized pitch and -duration of current melody note as inputs to dataset;
                add output note to list of uniques notes;
                add output note index as normalized output to dataset;

Since a melody track is stored as a sequence of notes, a note in the melody track and a note in the accompaniment track at the same index might not occur at the same time. Thus we iterate through the melody track and use the note closest in time in the accompaniment track for the algorithm.

Training LSTM and other recurrent neural networks is slow, as such the weights and structure of the network are saved to storage and is precomputed for each instrument and style of music.

16.2.1 *Results*

Since the network requires an input melody track to construct an accompaniment track it does not face the same problems such as lack of structure which most other non-recurrent neural networks do.

The output accompaniment can sometimes be pleasant, however a better timing relation is required between the input and output notes. Most of the produced melodies are repetitive, see figure 26. The network does not produce an accompaniment which is in harmony or in phase with the main melody track.

Figure 26: Accompaniment melody produced with a LSTM network

Overall the results of the LSTM network were average, in comparison with the Markov Chains. A melody generated by the LSTM network had better overall structure and theme but were repetitive. Markov chains produced better local melodies. The poor performance of the LSTM network is attributed to the implementation problems and the large training time required for good results.

ACCOMPANIMENT GENERATION WITH A FEED FORWARD NETWORK

17.1 INTRODUCTION

In this experiment an attempt is made to generate an accompaniment track for a melody using a feed forward network. One of the reasons for this experiment was to reduce the training time, as recurrent networks have a very large training time.

17.2 METHODOLOGY

For this prototype we construct a neural network with the following parameters:

1. all units use the sigmoid activation function $\frac{1}{e^{-x}}$
2. 1 input layer, 1 hidden layer, 1 output layer
3. 2 input units
4. 2 output units
5. 10 hidden units

The data for the network consists of the melody track and the accompaniment track. For each note in the melody track the note pitch and duration is given as inputs to the network, the output is the note pitch and duration of the corresponding note in the accompaniment track.

17.3 RESULTS

Figure 27 shows an accompaniment melody produced with the network. Since the network does not feature a recurrent topology there was no overall theme or structure. The network learns to output a certain note based on the input. The melodies produced had a better local structure than those produced by the recurrent network. The training speed and execution speed was significantly faster than that of the recurrent network.

Figure 27: Accompaniment melody produced with a LSTM network

ACCOMPANIMENT GENERATION WITH A MARKOV MODEL

18.1 INTRODUCTION

For this experiment an attempt is made to produce an accompaniment based on the frequency of certain notes occurring in the accompaniment track given the melody track.

More specifically, for each note in the melody track the odds are calculated for each possible accompaniment note that may occur at that time. Given an input melody, for each note in that melody an accompaniment note is produced according to the calculated probabilities. For this the topology of a HMM is utilized. A Markov Chain (without transition probabilities) is created and the emission probabilities are calculated.

18.2 METHODOLOGY

Algorithmus 3 : Constructing frequency table for model

Data : MIDI files that accord to a certain style;
Result : Emission frequency table for accompaniment notes;
for *each MIDI file* **do**

    Get main melody;

    for *each accompaniment track in file;*

    do

        for *each note* N_m, i *and previous note* $N_{m,i-1}$ *in main melody and corresponding note* N_a, i *in accompaniment melody;*

        do

            Increment frequency of N_a, i in frequency table for $\{N_m, i, N_m, i-1\}$;

Figure 28 illustrates a Markov Model. The states s_i denote the notes of the main melody (for a first order Markov Chain) and the states y_i denote the accompaniment notes for melody track. In the case of a HMM he states s_i are commonly hidden, however they are known in this case. The states y_i are the output notes that we are interested in generating.

For this experiment the states are setup up as a second order Markov Chain; that is the next note is determined by the previous two notes. For this model we are not interested in determining thee next state s_{i+1} thus the state transition matrix is not learnt. The emission probabilities of the output notes y_i are learnt instead.

Algorithmus 4 : Obtaining accompaniment melody

Data : Main melody sequence

Result : Accompaniment melody sequence

initialization;

Create empty accompaniment sequence;

for *each note* N_m, i *and previous note* $N_m, i-1$ *in main melody sequence* **do**

Lookup frequencies of possible output notes in frequency table for $\{N_m, i, N_m, i-1\}$;

Obtain probabilities for next notes in frequency table;

Obtain resulting note according to roulette selection;

Add resulting note to accompaniment sequence;

Return accompaniment sequence;

Figure 28: Illustration of a Hidden Markov Model

Once the emission probabilities are obtained, we can simply go through the states of the main melody and obtain the emission probabilities at each state. Roulette selection is performed based on the emission probabilities to determine a output note. The output note is then added to the accompaniment sequence. See algorithm 4

18.3 RESULTS

Figure 29: Accompaniment generated with Markov Model

Accompaniment generation with a Markov Model produces better results than with a LSTM network. The states of the model are the notes of the main melody and the observations are the notes of the accompaniment. For a short input melody sequence the resulting accompaniment sounds rather pleasant and fits with the main melody. For longer melodies the accompaniment diverges from the main melody and is not in harmony.

MELODY GENERATION WITH STOCHASTIC SAMPLING

19.1 INTRODUCTION

The idea here is to generate a melody from an extant music piece and modifying it to create a new melody.

Given some model λ we modify some piece p into a new piece p' such that $\Pr(p'|\lambda) > P(p|\lambda)$.

In order to modify the piece various sampling methods can be used. The idea is to replace a note of a piece with some other note in an attempt to increase the probability that the new piece come from the model.

19.2 METHODOLOGY

The model to be used will be a simple first order Markov chain. The model is able to evaluate the probability that a music piece or sequence of notes originated from it, i.e. it can calculate $\Pr(p|\lambda)$. The search space for new notes is the set of all notes contained within the Markov chain. The sampling to be performed is similar to Metropolis sampling. A random note is chosen from piece p and replaced with another note to produce p', the new note is kept if $\Pr(p'|\lambda) \geqslant \Pr(p|\lambda)$. The new piece will then be used in subsequent iterations.

Algorithmus 5 : Pseudocode for generating a melody using Stochastic Sampling

---

Data : Dataset that accords to a certain musical style
Result : Melody Sequence
model = MarkovChain();
Add all melody sequences in the dataset to the Markov Chain;
Let MaxIter be the maximum number of iterations;
iter = 0;
Let piece be a random melody sequence from the dataset;
while *iter < MaxIter* **do**

    iter = iter + 1;
    random_index = Random number between 0 and number of notes in piece;
    Let new_note be a random note in the search space;
    new_piece = piece.ReplaceNote(random_index, new_note);
    if *model.Probability(new_piece) >= model.Probability(piece)* **then**
        piece = new_piece;

Return piece;

Figure 30: Randomly selected input melody 1

Figure 31: Output melody using the stochastic sampling on melody 1

Figure 32: Randomly selected input melody 2

Figure 33: Output melody using the stochastic sampling on melody 2

Figure 30 shows the randomly selected melody used to generate the melody in figure 31 using stochastic sampling with replacement. The number of iterations for the first melodies were 100. Using a larger number of iterations of 200 more variety and change can be seen. Figure 32 shows the randomly selected melody used to generate the melody in figure 33 using the larger number of iterations.

19.4 CONCLUSION

The problem with this method is that it might not seem creative, even though it can produce interesting pieces. Another problem is that it might fall in a local minimum, where further replacements cannot improve the probability of the piece. Lastly since a first order Markov Model was used to calculate the probability of the piece it could not always produce changes which had a overall theme. Some changes felt out of place.

Overall the method produced pleasing results and indicates that modifications to existing pieces can be a efficient way to generate new pieces in a certain style.

In order to improve this method a different of more accurate model could be used other than the first order Markov chain. A HMM might yield better results.

Part V

DETAIL DESIGN

INTRODUCTION

In this part the design of the end user application will be discussed in detail. Topics covered include the final selection of algorithms in the application, the design of the user interface, a high level overview of the data structures used, the functionality and features of the application and finally some additional considerations will be reviewed.

TECHNOLOGY, PLATFORM AND PROGRAMMING LANGUAGE

21.1 INTRODUCTION

Some algorithms and features lend themselves to being implemented more easily in certain languages than other. A programming language needs to be chosen that does not restrain the functionality of the application. The four languages with most support include:

1. Python
2. C# (and other .NET languages)
3. Java
4. C++

Machine learning libraries exist for all of the above mentioned languages. Python has poor support for rich user interfaces. Java is a good contender and is cross platform. Haskell has strong expressiveness and powerful language features.

21.2 TRADE-OFF STUDY

Thus in order to make a final selection a trade off study was conducted. Additional factors included in the trade-off study were the expressiveness and amount of features in a language and the cross platform compatibility. These features have a lower weighing. Consideration is made to the end user and the user interface support is given the highest weighing.

Thus the factors considered are:

- Development speed - Estimated amount of programmer time required in each language.
- Available support and help - How easy it is to obtain help and support
- Available machine learning and pattern recognition libraries. Implementation of some algorithms prove difficult.
- User Interface (UI) design and libraries - Availability of libraries to design a good user interface and experience
- Cross platform compatibility - How well does the language, features and libraries run on other platforms
- Expressiveness and features - Language features which ease development. Expressive languages facilitate advanced data structures and algorithms in few Lines of Code (LOC)

Table 5: Programming languages trade-off study

	Weight	C#	Java	Python	C++	Haskell
Development speed	0.5	8	7	8	5	9
Available support and help	0.2	8	8	8	6	4
Machine learning libraries	0.4	6	8	9	8	4
UI design and libraries	0.6	10	8	5	7	3
Cross platform compatibility	0.2	6	10	10	10	10
Expressiveness and features	0.4	7	6	8	5	10
Total	2.3	18	17.5	17.4	15.1	14.7

21.3 DISCUSSION

Winner: *C#*[1]

.NET languages facilitate the use of WPF - a graphical system for rendering user interfaces, which will allow a rich user experience to be developed. Java and Python are good contenders.

C# additionally has the following advantages:

1. Ability to easily create rich user interfaces.

2. C# runs on the CLR and allows access to the large .NET library[2]

3. Powerful language constructs such as operator overloading, event handling, delegates, powerful multi-threading support and generics.

---

1 The algorithmic code is cross platform and can be run through Mono, even though the graphical front end developed in Windows Presentation Foundation (WPF) cannot.
2 Additionally the IKVM project allows access to Java libraries

22

SELECTION OF ALGORITHMS

22.1 INTRODUCTION

In this chapter the selected algorithms for the end user application are chosen. The algorithms for consideration are:

1. Melody generation with Genetic Algorithms.
2. Melody generation with Hidden Markov Models.
3. Accompaniment generation with Hidden Markov Models.
4. Accompaniment generation with a LSTM network.
5. Accompaniment generation with a feed forward ANN.
6. Instrumental generation with Markov Chains.

The selection of algorithms will be done by means of elimination. Algorithms that yielded sub minimum results will not be selected for inclusion into the end user application.

22.2 TRADE-OFF STUDIES

For the final end-user application only a subset of algorithms will be chosen. In order to quantitatively select good algorithms a trade-off study will be done on the investigated algorithms. The factors considered are:

- Training time - The amount of time taken to train the algorithm on initial data.
- Generation time - The amount of time the algorithm takes to produce a melody after it is trained.
- Flexibility - How well the algorithm can accommodate different input data. How easily it can be extended or modified.
- Pleasantness - Subjective measure of good a melody is.

Table 6 shows the trade-off study for algorithms for melody generation. Melody generation is simply the generation of a single track consisting of a list of notes. From the trade-off study we can see that the HMM and Feed Forward Neural Network (FFANN) performed poorly and thus will not be selected.

Table 7 indicates the trade-off study for algorithms for accompaniment generation. Accompaniment generation is a list of notes that have to accompany the main melody. From the trade-off study we can see that the HMM and LSTM performed poorly and thus will be excluded from the main application.

Table 6: Trade-off study for algorithms for melody generation

Melody generation					
	Weight	GA	FFANN	MC	HMM
Training time	0.3	10	4	8	5
Generation time	0.6	4	7	8	7
Flexibility	0.6	10	7	6	6
Pleasantness	0.7	5	4	7	5
Total	2.2	14.9	12.4	15.7	12.8

Table 7: Trade off study for algorithms for accompaniment generation

Accompaniment generation					
	Weight	HMM	FFANN	MM	LSTM
Training time	0.3	6	4	8	3
Generation time	0.6	7	7	8	6
Flexibility	0.6	6	7	6	5
Pleasantness	0.7	3	5	6	3
Total	2.2	11.7	13.1	15	9.6

Table 8: Comparison of frequency metric time and pleasantness

	Time (ms)	Pleasantness
Chromatic Tone	8992	4
Chromatic Tone Distance	11017	5
Chromatic Tone Duration	13688	5
Melodic Bigram	3900	5
Melodic Interval	9037	5
Pitch	14001	4
Pitch Distance	-	-
Rhythm	367	4
Rhythmic Bigram	4729	7
Rhythmic Interval	5501	6

For the frequency metrics to be used with the cosine similarity fitness function we refer to table 3. The pleasantness and time for each metric is repeated in table 8 for convenience. The Pitch Distance metric was excluded as it took too long. The chromatic tone, pitch and rhythm metrics will also be excluded as all scored low pleasantness.

22.3 CONCLUSION

From the trade-off study the selected algorithms were chosen. In summary, the selected algorithms for the end-user application were:

1. Genetic Algorithms for melody generation.
2. Markov Chains for melody generation (with percussion support).
3. Feed Forward Artificial Neural Network for accompaniment generation.
4. Markov Model for accompaniment generation.

Only the cosine similarity fitness function will be used, as the NCD fitness function performed poorly. The metrics for the cosine similarity fitness function were selected as:

1. Chromatic Tone Distance
2. Chromatic Tone Duration
3. Melodic Bigram
4. Melodic Interval
5. Rhythmic Bigram
6. Rhythmic Interval

Using trade-off studies the worst results were eliminated from all algorithms.

FEATURES AND FUNCTIONALITY

The main goal of the application is to generate melodies according to a certain styles. The user is to select a certain category of music and the application will generate a musical piece in that style. The premise is to generate the melodies in a certain style using machine learning techniques that will use the data present in the corpus of MIDI files.

Since training machine learning algorithms take up a lot of time it is not feasible to train the algorithm on the fly. Training will take place beforehand and the resulting data structure of the algorithm will be cached to a file. For example a Neural Network will save its topology and weights to a file; this allows the application to load the file, present inputs to the neural network and output a result without spending time learning the optimal weights for the network.

Additional features would be to allow the user to save a generated melody to a file (in MIDI or a custom file format), load and play an existing MIDI file (for comparison or other aims), re-randomize a generated composition and view more information or properties of the generated music piece.

The application is to be used by end users which do not have background on machine learning. The application is to be as user-friendly as possible while still having flexibility in composing songs. The user interface is to visually display the song and playback of it. Playback functionality such as playing and stopping a song are mandatory.

To summarize, the application will:

1. Implement the selected algorithms and allow them to be used for song generation.
2. Allow saving and loading of MIDI files
3. Generate a monophonic melody according to a certain style. Allow the user to select which algorithm to use and adjust high level parameters
4. Playback functionality including playing, pausing and stop a song.
5. Feature a user friendly interface.
6. Visually display the generated melody.
7. Allow the user to randomize the song based on the selected settings.
8. Cache machine learning data to improve performance of the application.

GRAPHICAL USER INTERFACE

24.1 INTRODUCTION

In this chapter the design of the graphical user interface is presented, along with the features and functionality of the application.

As mentioned above the user interface is to have the following features

- Visually display the generated song.
- Selecting the category of melodies to be generated.
- Select algorithm and high level parameters.
- Buttons to allow for saving, loading, pausing, pausing, stopping and playing of songs.

The user interface will be developed using WPF and Extensible Markup Language (XAML).

24.2 PLAYBACK

Figure 34 indicates the starting interface of the application. This interface provides controls for playback functionality, is responsible for displaying a visual of the song being currently played, provides saving and loading functionality and allows the user to randomize the current song using the parameters set in the compose view.

24.3 COMPOSITION

The interface as shown in figure 35 allows the user to select the music style or category for music to be generated in. A track is a sequence of notes played in a particular instrument. A user is able to add new tracks using the provided button, which takes them to a new window as in figure 36. If a user double clicks on a specific track the frequency metrics of the track are shown as in figure 37.

24.4 TRACK GENERATION

Figure 36 allows the user to generate a track using a specific technique. This interface allows the user to select the type of technique for track generation and also displays the progress. Once a track is generated it can be added to the list of tracks as in figure 35.
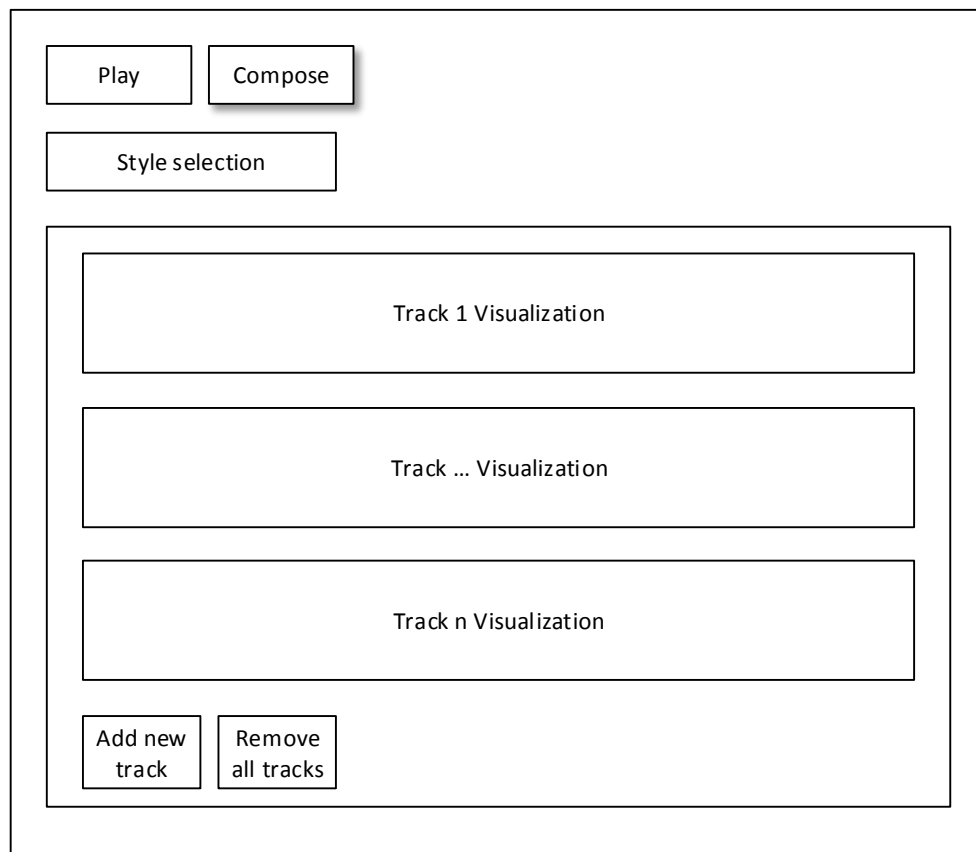
Figure 34: User interface for playback functionality

Figure 35: User interface for composition functionality

Genetic	Instrumental	Accompaniment	Random	Load

Options

| Generate	Progress indicator	Add	Close

Figure 36: User interface for melody generation

Figure 37: User interface for displaying metrics of a melody

24.5 STATISTICS

Figure 37 displays a list of frequency metrics about the track selected. An example metric is the frequency of musical intervals within a melody (mi = $|p_i - p_{i-1}|$, where p_i is the pitch at index i); more information about frequency metrics can be found in section .

24.6 VISUALIZATION

The visualization blocks in figures 35 and 36 will be a simple visual displaying the notes over time.

AlphaTab is a cross platform music notation and guitar tablature rendering library. It is the only music notation library available on C# which will fit the need for visualization of notes required. Figure 38 indicates a music sheet rendered with Alphatab. Note that Alphatab is geared towards guitarists and provides a lot more information which is unwanted for a simple visualization.

Since the library is open source it will be modified to suit our needs. Guitar tabs, timing information, lyrics and any other excessive information will be

Figure 38: Example of a music sheet rendered with alphatab

Figure 39: Desired notes visualization

removed. In addition it is necessary to show which note is currently played when a melody is played. The active note will be displayed in red.

Figure 39 displays the designed simple notes visualization required for the application. Alphatab will be modified in order to produce a result similar to this.

24.7 STATE CONTROL

Figure 40 shows the state diagram of the user interface. From the figure it can be seen that the Playback UI controls playback of a single composition. From the composition UI the necessary tools and functionality can be accessed to create new melodies.
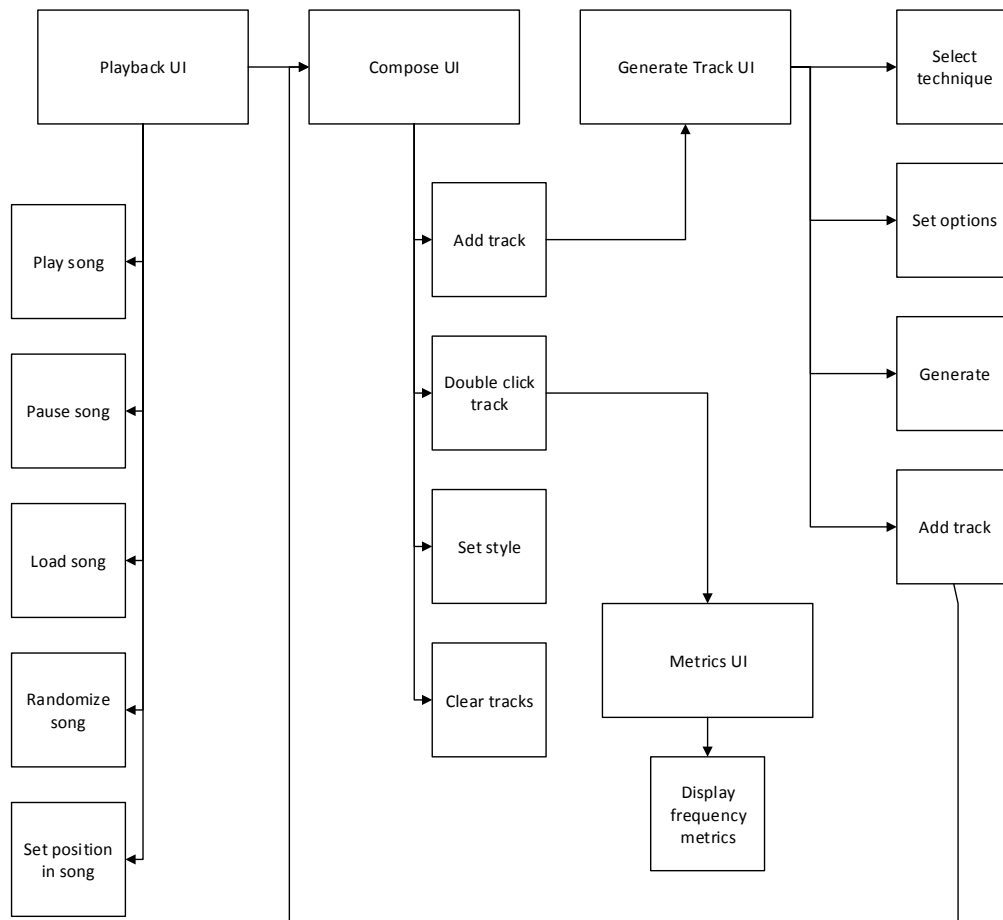
Figure 40: State diagram of the user interface

FINAL CONSIDERATIONS

Part VI

DENOUEMENT

RESULTS

26.1 INTRODUCTION

This chapter outlines the results, including the design of the user interface, functionality of the application and a short summary of the results of the incorporated algorithms. The majority of the results have already been covered in part iv, however a discussion is provided.

26.2 USER INTERFACE

Figure 41 shows the frequency metrics for the first track of the Harry Potter opening theme song.

Figure 42 shows the track generator with all of the available tabs. These tabs provide the following features:

1. Genetic - Generates a melody using genetic algorithm. Allows selection of which frequency metrics to use with the cosine similarity fitness function.

2. Instrumental - Generates a melody using Markov Chains for a particular instrument.

3. Accompaniment - Generates a accompaniment melody for a main melody track using a Markov model or a ANN.

4. Load - Loads an existing melody or modifies an extant melody using stochastic sampling

5. Random - Generates a random melody using a random walk in a specific scale.

Figure 43 shows the play tab of the main window. The play tab offers access to all playback functionality including playing, stopping and seeking the position of a melody. In addition the play tab allows the user to re-randomize the melody. The current melody can be saved or an existing melody can be loaded. Lastly the play tab shows a visualization of the current song playing. A note being played is highlighted in red.

Figure 44 shows composition tab of the main window. The visualization shows the current composition. A new track can be added to the composition on which the track generator window is shown. In addition a context menu provides the following features:

- Modification of the tempo of a track.
- Transposition of all note pitches on a track.
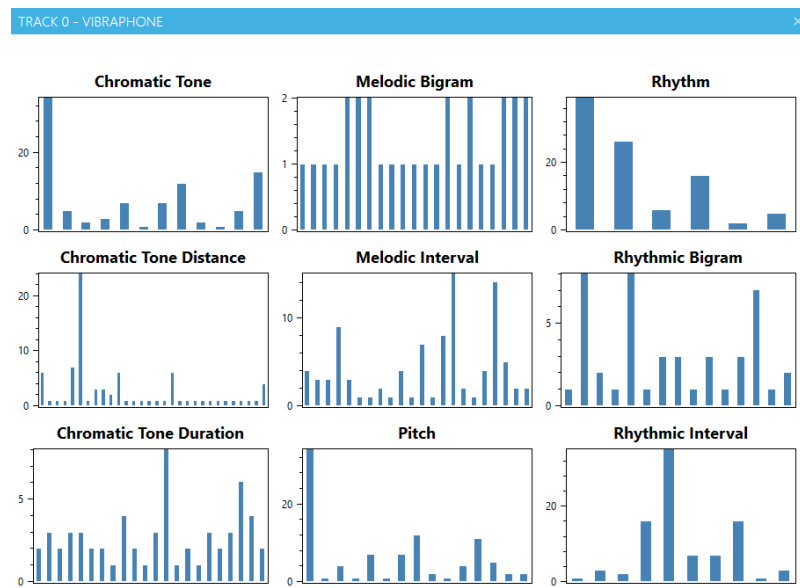- Ability to show the frequency metrics of the current track.

Figure 41: Frequency metrics for the first track Harry Potter Opening Theme Song
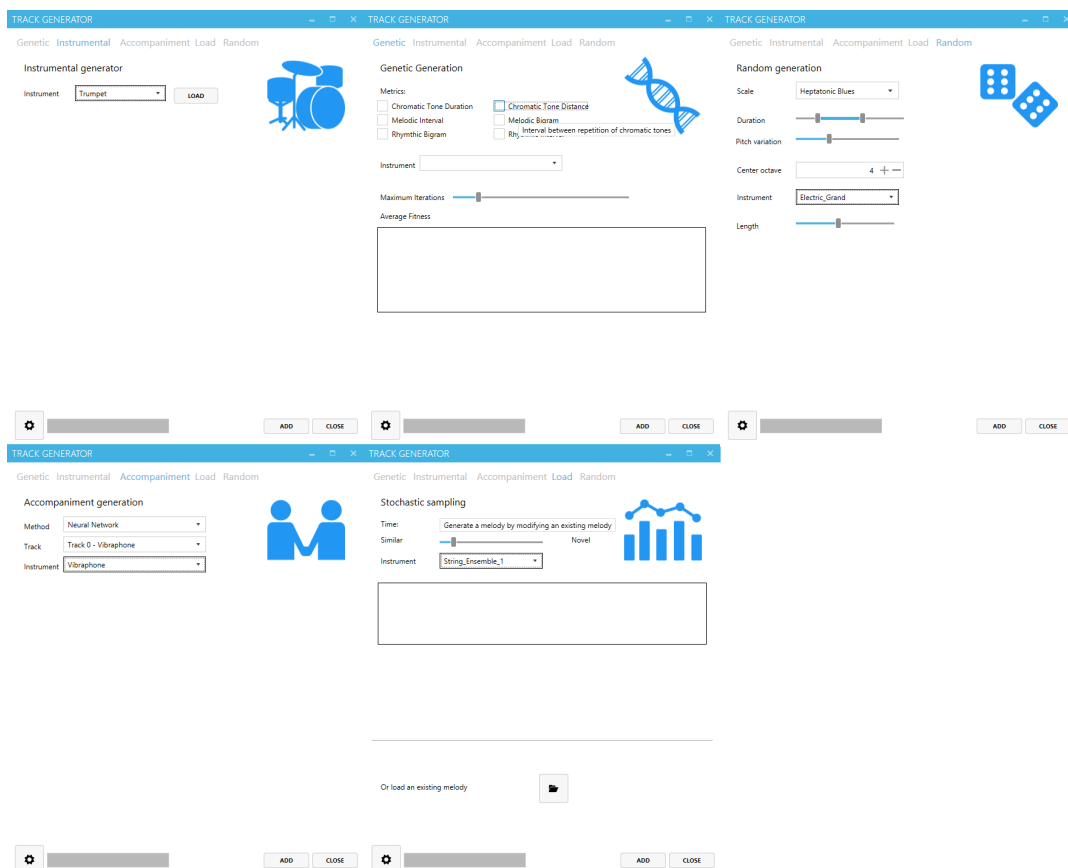
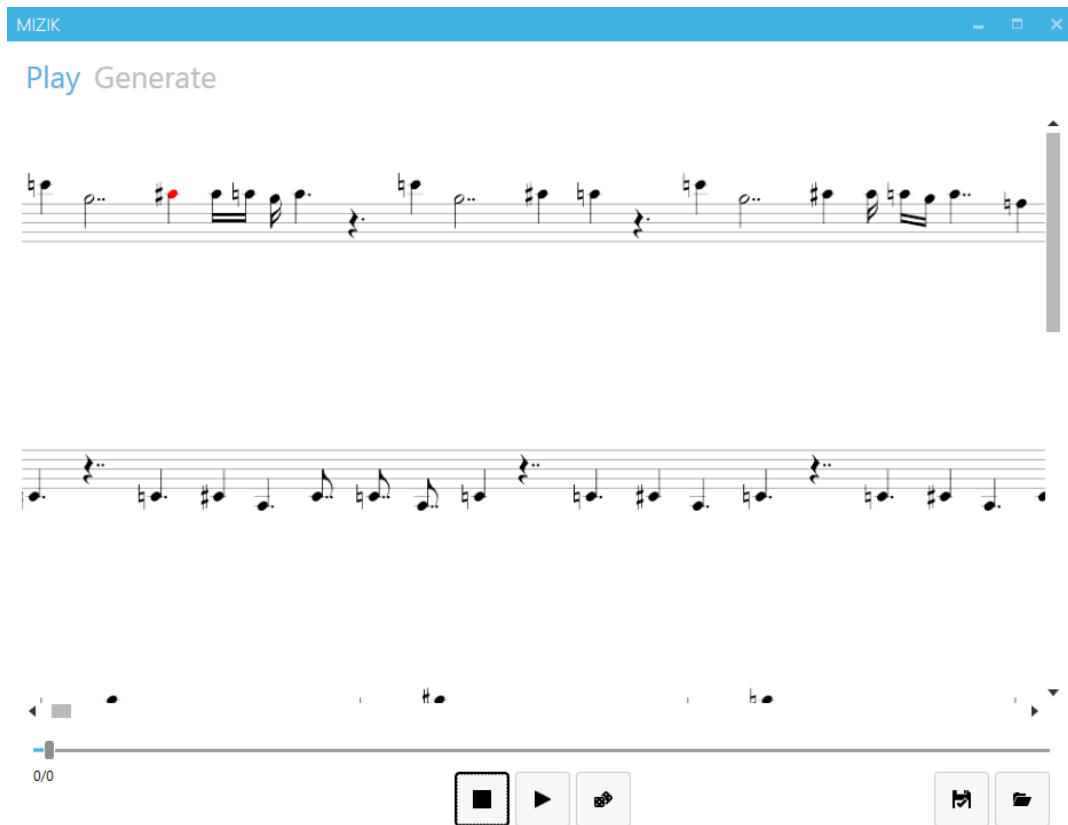

Figure 42: Tabs of the track generator window

Figure 43: Playback tab of main window

Figure 44: Playback tab of main window

Figure 45: Excerpt of a melody with acoustic bass and percussion

Figure 46: Excerpt of a melody with a string ensemble

As can be seen from the figures the user interface features a clean, simple and friendly design that is based on the Metro style in Windows.

26.3 ALGORITHMS

As has been previously mentioned the results of the algorithms were discussed in part iv. Some modification and changes were made. Most significantly the parameters of the algorithms can be modified by using the track generator. This provides some flexibility to the user.

In addition a random walk algorithm was implemented. This algorithm randomly walks through the selected scale. Although the algorithm does not compose music in a certain style and is quite simple and repetitive it does provide a useful comparison and can be used to generate simple melodies.

Some interesting melodies were produced throughout. Figure 45 shows an excerpt of a melody containing acoustic bass and percussion instruments which play well together. Figure 46 is a melody with a string ensemble which has a deep and spacey feel. Figure 47 is a melody containing a trumpet and drums in a Jazz vibe.

Overall some instruments did better than others. Especially with the Markov models, however this might be due to the input data.

Some interesting subjective observations made include:

1. The trumpet in the Jazz style performed well in most algorithms

2. The violin performed poorly

3. Jazz generation all round produced more pleasant melodies

4. Classical produced average results. Classical contained a mixture of classical music. It would be interesting to see how the algorithms performed on a subset of classical music, for example on Bach's music

Figure 47: Excerpt of a melody with a trumpet and drum

5. Video game music produced interesting music especially background music, for example for use in video games, or elevators or similar applications.

6. More work is required in accompaniment generation. For longer pieces the tracks desynchronize.

CONCLUSION

The aim of this project was to design an application that is able to compose melodies in a certain musical style using machine learning algorithms. A brief survey was done on existing attempts to compose melodies using machine learning techniques.

Evolutionary algorithms are versatile and are well suited to the the task of musical composition. The task of designing a good fitness function for a genetic algorithm is rather formidable. Musical composition is a emotional human process and is subjective in nature. This makes formalization of what constitutes a good melody difficult. An attempt was made by using the NCD to rate the fitness of melodies, however the results were not satisfactory. Decent results were obtained by using frequency metrics in conjunction with a similarity measure. This approach provides high fitness to individuals which have similar statistical properties to the target piece. The problem with the approach is that a high statistical similarity may be achieved but this does not necessarily imply that a individual will be perceived as pleasant or good.

A different approach was taken thereafter, other than relying on statistical properties and requiring a function to quantify the pleasantness of a melody. A probabilistic approach was taken to predict the pattern of melodies. By employing the Markov property - the probability distribution of future states depend only upon the present states simplifications can be made and Markov Models can be designed. The simplest Markov Model, the Markov Chain produced pleasant results. A third order Markov Chain, where the next note is determined by the preceding 3 notes produced pleasant results without sacrificing too much novelty and avoiding too much similarity. By employing another Markov Model, attempt was made to produce an accompaniment for a melody, by calculating the probabilities for an accompaniment note by the preceding 2 notes in the melody track. This Markov Model also yielded satisfactory results however more work is required on timing and synchronization. Overall these probabilistic models yield good results, however the melodies produce lack an overarching theme and structure. Furthermore since Markov Models require a input dataset in order to calculate the state transition and observation emission probabilities, the novelty and composition value are low. These models produce pleasant results but are not creative or original. This is in contrast to the evolutionary algorithms.

Further works is needed to develop an algorithm that is able to compose melodies that have an overarching theme or structure. Further work can be done using recurrent neural networks in a variety of topologies or using

a custom HMM that is designed for a single musical style. These two techniques were used in a simplistic way to generate the accompaniment to a melody with acceptable results, however these models are powerful and further work can be done.

The results of most the algorithms could be improved by incorporating domain knowledge into the system. Musical knowledge could constrain the algorithm to produce more acceptable results. A hybrid system that incorporates expert knowledge may reduce the number of unpleasant patterns produced and improve note transitions and so forth.

The application used a number of simplifications. A major simplification was that only monophonic melodies were considered, no chords. A more advanced system can be designed that applies the investigated algorithms to chords as well. Further algorithms could be used for harmonization.

The application conforms to the specifications and is able to satisfactorily produce melodies in Classical, Jazz, Drum and Video Games styles using Markov Models, Neural Networks and Genetic Algorithms.

28

FUTURE WORK

In this chapter some ideas for future work on the application are summarized. The ideas for future work include:

- If the application is to be used for background application let the algorithms generate piece of infinite length and continuously play the melodies.
- Further work is required on accompaniment generation.
- Classical music can be split into multiple sub categories. One idea is to have a different style for each famous composer.
- Support for chords and polyphonic melodies. Algorithms can be extended to support this.
- Further work can be done using recurrent neural networks for melody generation rather than accompaniment generation.
- The algorithmic music composition field is large. There are a multitude of other algorithms which can be implemented.
- Better datasets can be used. Each individual MIDI file can be reviewed for inclusion.

Part VII

APPENDIX

A

ALGORITHMS

A.1 BACKPROPOGATION

For a feed forward network[1] with n_{in} inputs, n_{hidden} hidden units and n_{out} output units the back-progogation algorithm works as follows

1. Initialize network with random weights
2. Repeat until algorithm terminates
3.
   - $\forall (\vec{x}, \vec{t}) \in$ training examples do
     - a) Input instance $\vec{x}$ to network and compute $o_u \ \forall u \ \in$ network
     - b) For each network output unit k calculate error term
       $$\delta_k \leftarrow o_k(1 - o_k)(t_k - o_k)$$
     - c) For each hidden unit h calculate the error term δ_h
       $$\delta_h \leftarrow o_h(1 - o_h) \sum_{k \in outputs} w_{kh} \delta_k$$
     - d) Update each network weight w_{ji}
       $$w_{ji} \leftarrow w_{ji} + \eta \delta_j x_{ji}$$

The complete derivation of the back-progogation algorithm for feed forward artificial neural networks can be found in Mitchell (2007) [44]

---

1 Other topologies exist with more complex methods for obtaining the weights

PLATFORM SPECIFICATIONS

Table 9: System specifications

Operating System	Windows 8.1
Memory	8GB 1333Mhz DDR3
CPU	AMD X3 450 3.2GHZ 4 cores

Table 10: Development specifications

Language	C#
IDE	Visual Studio 2015 Community
.NET version	4.5

All tests were run with the specifications as given in table 9. The development environment specifications are given in figure 10.

SOURCE CODE

Listed here is the framework used to generate the results mentioned in this project.

C.1 DOTNETMUSIC

DotNetMusic contains the core music/MIDI functionality and includes:

- Reading and Writing to MIDI files using NAudio
- Core data structures for music representation
- Core playback functionality
- WPF control for displaying notes on a music sheet
- Saving and loading of core music data structures using Google protocol buffers

DotNetMusic can be found at github.com/stefan-j/DotNetMusic

C.2 DOTNETLEARN

Some additional machine learning implementation and useful math features.

- Implementation of Markov Model
- Implementation of Hidden Markov Model with multiple observation Baum Welch Learning
- Codification of data
- Other statistical features

DotNetLearn can be found at github.com/stefan-j/DotNetLearn

C.3 GENETICMIDI

GeneticMIDI contains the core functionality for algorithmically generating music.

- Generators for producing music algorithmically using Markov Models, Markov Chains, Neural Networks.
- Genetic Algorithms (Genetic Program (GP) tree structures)
- Fitness functions for Genetic Algorithms including NCD, Cosine Similarity (and frequency metrics)
- Random generation of notes

- The developed front-end visualizer application for producing music in a user-friendly manner

Libraries used include:

- A-Forge - For the neural networks and evolution of GAs.
- NAudio - For reading, writing and playing of MIDI data.
- IKVM - For the **LTSM!** (**LTSM!**) network.
- Proto.Net - .Net implementation of Google protocol buffers for saving of data structures

GeneticMIDI can be found at github.com/stefan-j/GeneticMIDI

[1] D. Eck and J. Schmidhuber, "A first look at music composition using lstm recurrent neural networks," *Istituto Dalle Molle Di Studi Sull Intelligenza . . .*, 2002.

[2] P. M. Gibson and J. A. Byrne, "Neurogen, musical composition using genetic algorithms and cooperating neural networks," in *Artificial Neural Networks, 1991., Second International Conference on*, pp. 309–313, IET, 1991.

[3] M. C. Mozer, "Neural network music composition by prediction: Exploring the benefits of psychoacoustic constraints and multi-scale processing," *Connection Science*, vol. 6, no. 2-3, pp. 247–280, 1994.

[4] J. a. Biles, "GenJam: A genetic algorithm for generating jazz solos," *Proceedings of the International Computer Music Conference*, pp. 131–131, 1994.

[5] A. Sorensen, "Music Composition in Java."

[6] C.-C. Chen and R. Miikkulainen, "Creating melodies with evolving recurrent neural networks," *IJCNN'01. International Joint Conference on Neural Networks. Proceedings (Cat. No.01CH37222)*, vol. 3, 2001.

[7] M. Alfonseca, M. Cebrian, and a. Ortega, "A Fitness Function for Computer-Generated Music using Genetic Algorithms," *WSEAS Transactions on Computers*, 2006.

[8] J. H. Jensen, *Evolutionary Music Composition.* PhD thesis, 2013.

[9] J. a. Biles, P. G. Anderson, and L. W. Loggi, "Neural Network Fitness Functions for a Musical IGA," *Architecture*, 1996.

[10] A. R. Burton and T. Vladimirova, "Genetic Algorithm Utilising Neural Network Fitness Evaluation for Musical Composition," in *In Proceedings of the 1997 International Conference on Artificial Neural Networks and Genetic Algorithms*, pp. 220–224, Springer-Verlag, 1997.

[11] B. Manaris, P. Roos, P. Machado, D. Krehbiel, L. Pellicoro, and J. Romero, "A corpus-based hybrid approach to music analysis and composition," *Proceedings of the National Conference on Artificial Intelligence*, vol. 22, p. 7, 2007.

[12] B. Manaris, P. Machado, C. Mccauley, J. Romero, and D. Krehbiel, "Developing Fitness Functions for Pleasant Music : Zipf's Law and Interactive Evolution Systems," *EvoWorkshops*, pp. 498–507, 2005.

[13] M. Unehara and T. Onisawa, "Construction of Music Composition System with Interactive Genetic Algorithm," *Evaluation*.

[14] L. Spector and A. Alpern, "Induction and Recapitulation of Deep Musical Structure," in *In Proceedings of the IJCAI-95 Workshop on Artificial Intelligence and Music*, pp. 41–48.

[15] B. Johanson and R. Poli, "{GP}-Music: An Interactive Genetic Programming System for Music Generation with Automated Fitness Raters," *Genetic Programming 1998: Proceedings of the Third Annual Conference*, pp. 181–186, 1998.

[16] K. McAlpine, E. Miranda, and S. Hoggar, "Making Music with Algorithms: A Case-Study System," *Computer Music Journal*, vol. 23, pp. 19–30, 1999.

[17] M. Farbood and B. Schoner, "Analysis and synthesis of Palestrina-style counterpoint using Markov chains," *Proceedings of the International Computer Music Conference*, pp. 471–474, 2001.

[18] M. Allan and C. K. I. Williams, "Harmonising Chorales by Probabilistic Inference," *Advances in Neural Information Processing Systems*, p. 2004, 2004.

[19] M. Ling and R. Sharp, "Melody classification using a similarity metric based on kolmogorov complexity," *Conference on Sound and Music Computing*, 2004.

[20] MIDI Manufacturers Assosciation, "MIDI 1.0 Detailed Specification," first edition of a recommendation, MIDI Manufacturers Assosciation, 1995.

[21] B. L. Jacob, "Composing with genetic algorithms," *In Proceedings of the 1995 International Computer Music Conference*, no. September, pp. 452–455, 1995.

[22] W. Langdon, *A Field Guide to Genetic Programing (Summary for Wyvern)*. No. March, 2008.

[23] M. Minsky and O. Laske, "a C Onversation With M Arvin M Insky," *AI Magazine*, vol. 13, no. 3, pp. 31–45, 1992.

[24] L. E. Baum, T. Petrie, G. Soules, and N. Weiss, "A Maximization Technique Occurring in the Statistical Analysis of Probabilistic Functions of Markov Chains," *The Annals of Mathematical Statistics*, vol. 41, no. 1, pp. 164–171, 1970.

[25] H. Jarvelainen, "Algorithmic musical composition," *Seminar on content creation, Telecommunications software and multimedia laboratory, Helsinki University of Technology*, p. 11, 2000.

[26] D. Conklin, "Music Generation from Statistical Models," *Proceedings of the AISB 2003 Symposium on Artificial Intelligence and Creativity in the Arts and Sciences*, pp. 30–35, 2003.

[27] C. M. Bishop and Others, "Neural networks for pattern recognition," 1995.

[28] B. Laden and D. H. Keefe, "The representation of pitch in a neural net model of chord classification," *Computer Music Journal*, pp. 12–26, 1989.

[29] G. G. Wiggins, G. Papadopoulos, S. Phon-amnuaisuk, and A. Tuson, "Evolutionary methods for musical composition," *In International Journal of Computing Anticipatory Systems*, 1998.

[30] D. Matić, "A genetic algorithm for composing music," *Yugoslav Journal of Operations Research*, vol. 20, no. 1, pp. 157–177, 2010.

[31] M. Y. Lo, *Evolving Cellular Automata for Music Composition with Trainable Fitness Functions*. PhD thesis, 2012.

[32] A. Horner and D. E. Goldberg, "Genetic Algorithms and Computer-Assisted Music Composition," 1991.

[33] R. Cilibrasi and P. M. B. Vitanyi, "Clustering by compression," *Information Theory, IEEE Transactions on*, vol. 51, pp. 1523–1545, Apr. 2005.

[34] N. Tokui and H. Iba, "Music composition with interactive evolutionary computation," *Proceedings of the Third International Conference on Generative Art*, pp. 215–226, 2000.

[35] M. Alfonseca, M. Cebrián, and A. Ortega, "A simple genetic algorithm for music generation by means of algorithmic information theory," *2007 IEEE Congress on Evolutionary Computation, CEC 2007*, pp. 3035–3042, 2007.

[36] A. N. Kolmogorov, "On tables of random numbers," *Theoretical Computer Science*, vol. 207, no. 2, pp. 387–395, 1998.

[37] C. H. Bennett, P. Gacs, M. Li, P. M. B. Vitanyi, and W. H. Zurek, "Information distance," *Information Theory, IEEE Transactions on*, vol. 44, pp. 1407–1423, July 1998.

[38] M. Li, X. Chen, X. Li, B. Ma, and P. M. B. Vitanyi, "The similarity metric," *Information Theory, IEEE Transactions on*, vol. 50, pp. 3250–3264, Dec. 2004.

[39] Z. Cataltepe, A. Sonmez, and E. Adali, "Music Classification Using Kolmogorov Distance," *Representation in Music/Musical Representation Congress*, no. Bishop 57, 2005.

[40] M. Alfonseca, M. Cebrián, and A. Ortega, "Evolving computer-generated music by means of the normalized compression distance," in *WSEAS Transactions on Information Science and Applications*, vol. 2 of *SMO'05*, (Stevens Point, Wisconsin, USA), pp. 1367–1372, World Scientific and Engineering Academy and Society (WSEAS), 2005.

[41] C. Mckay, "Automatic Genre Classification of MIDI Recordings," *Technology*, vol. MA, no. June, pp. 1–150, 2004.

[42] G. K. Zipf, *Human behavior and the principle of least effort: an introduction to human ecology*. Addison-Wesley Press, 1949.

[43] G. A. Carpenter and S. Grossberg, *Adaptive resonance theory*. Springer, 2010.

[44] T. M. Mitchell, *Machine Learning*, vol. 4 of *McGraw-Hill Series in Computer Science*. McGraw-Hill, 1997.